

CS 110

Computer Architecture

Single-Cycle CPU

Datapath & Control

Instructor:
Sören Schwertfeger

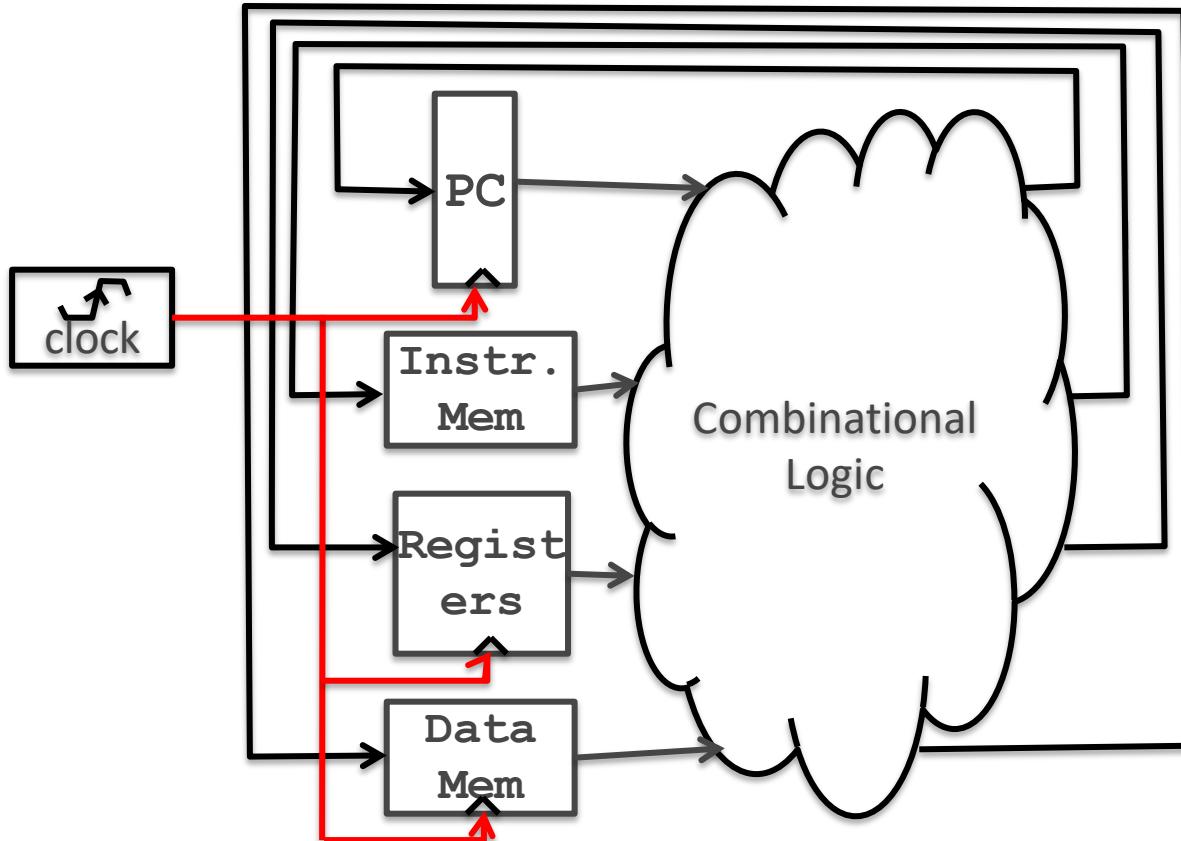
<http://shtech.org/courses/ca/>

School of Information Science and Technology SIST

ShanghaiTech University

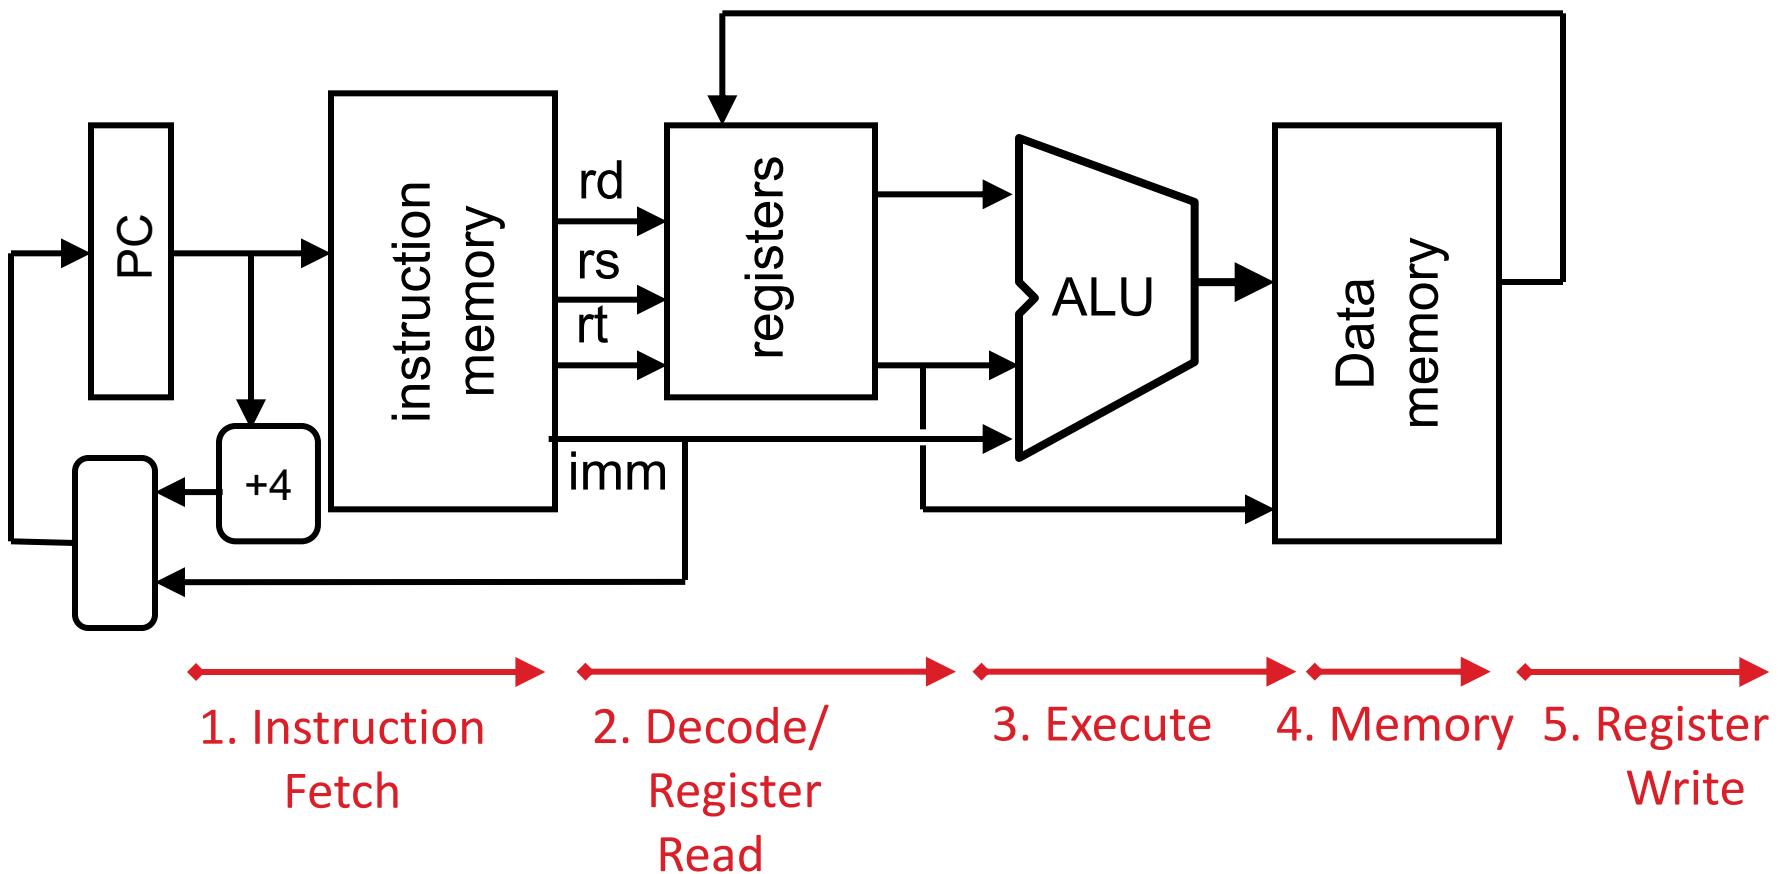
Slides based on UC Berkley's CS61C

One-Instruction-Per-Cycle RISC-V Machine



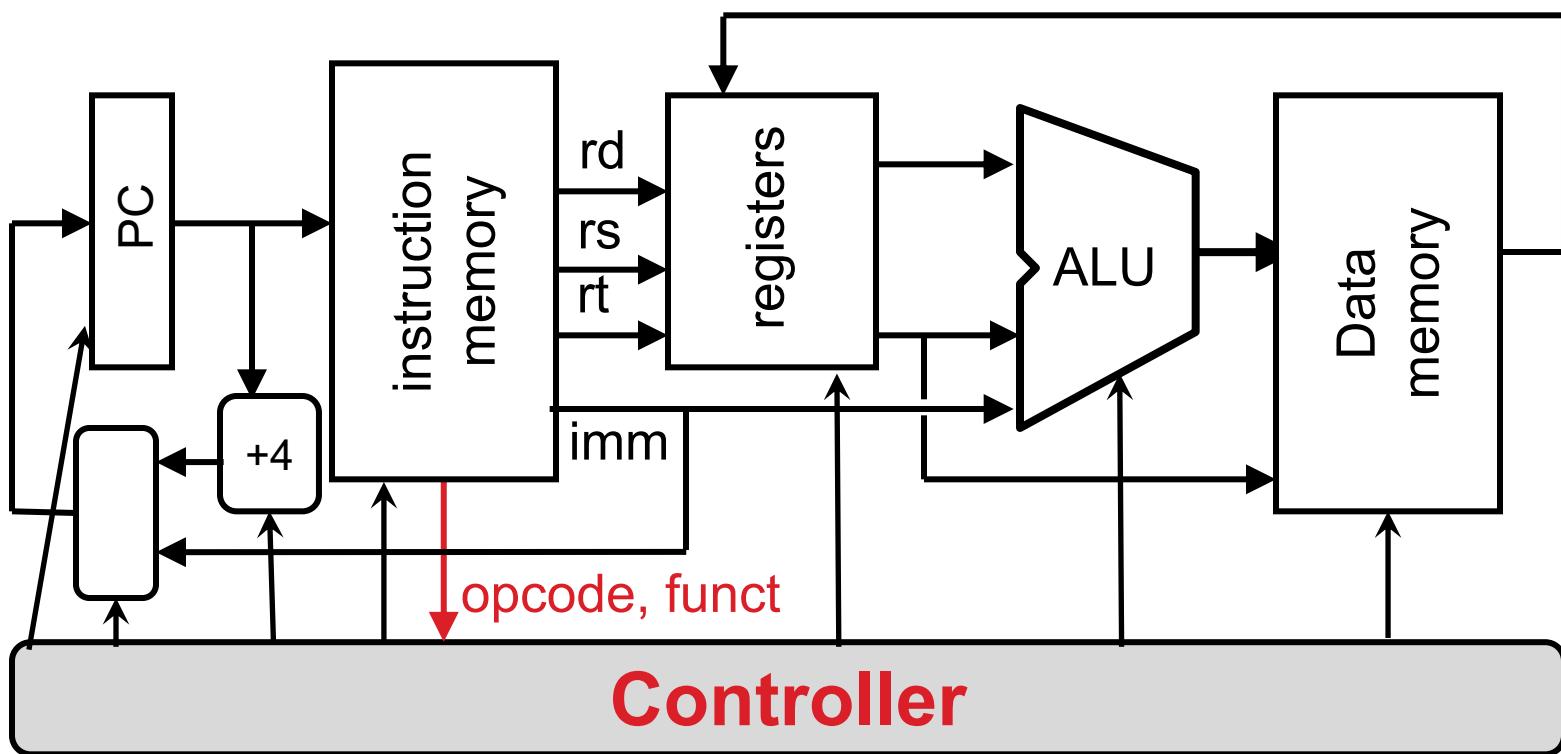
- One clock tick => one instruction
- Current state outputs => inputs to combinational logic => outputs settle at the values of state before next clock edge
- Rising clock edge:
 - all state elements are updated with combinational logic outputs
 - execution moves to next clock cycle

Stages of Execution on Datapath



Datapath and Control

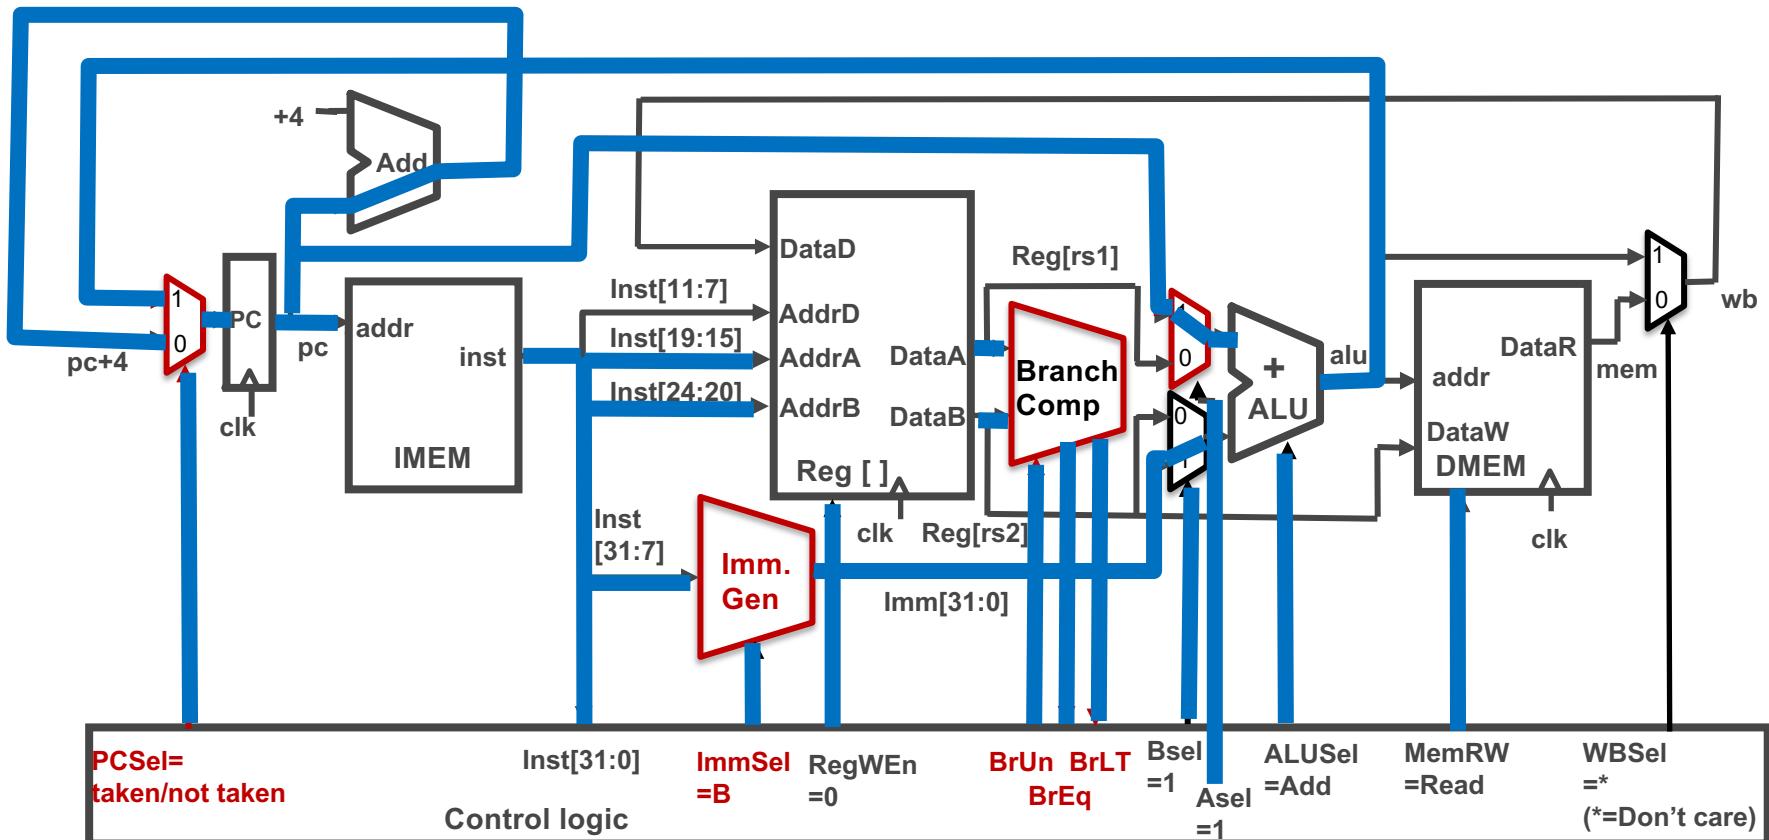
- Datapath designed to support data transfers required by instructions
- Controller causes correct transfers to happen



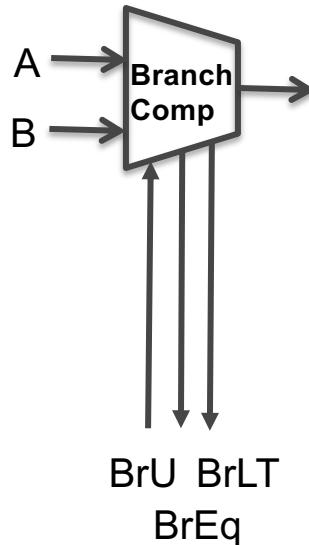
Branches

- Different change to the state:
 - $\text{PC} = \begin{cases} \text{PC} + 4, & \text{branch not taken} \\ \text{PC} + \text{immediate}, & \text{branch taken} \end{cases}$
- Six branch instructions:
BEQ, BNE, BLT, BGE, BLTU, BGEU
- Need to compute $\text{PC} + \text{immediate}$ and to compare values of rs1 and rs2
 - But have only one ALU – need more hardware

Adding Branches



Branch Comparator

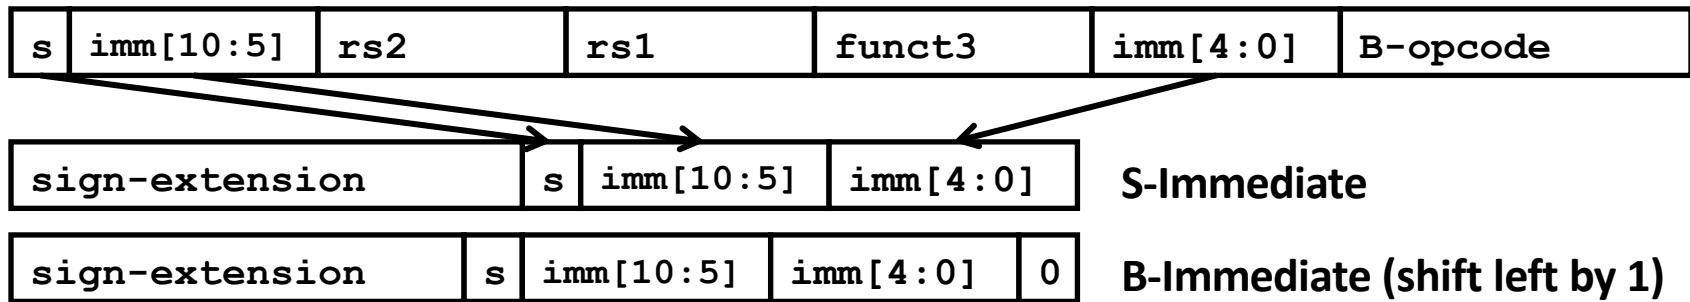


- $\text{BrEq} = 1$, if $A=B$
- $\text{BrLT} = 1$, if $A < B$
- $\text{BrUn} = 1$ selects unsigned comparison for BrLT , $0=\text{signed}$
- BGE branch: $A \geq B$, if $\overline{\overline{A < B}} = !(A < B)$

$$\overline{\overline{A < B}} = !(A < B)$$

Branch Immediates (In Other ISAs)

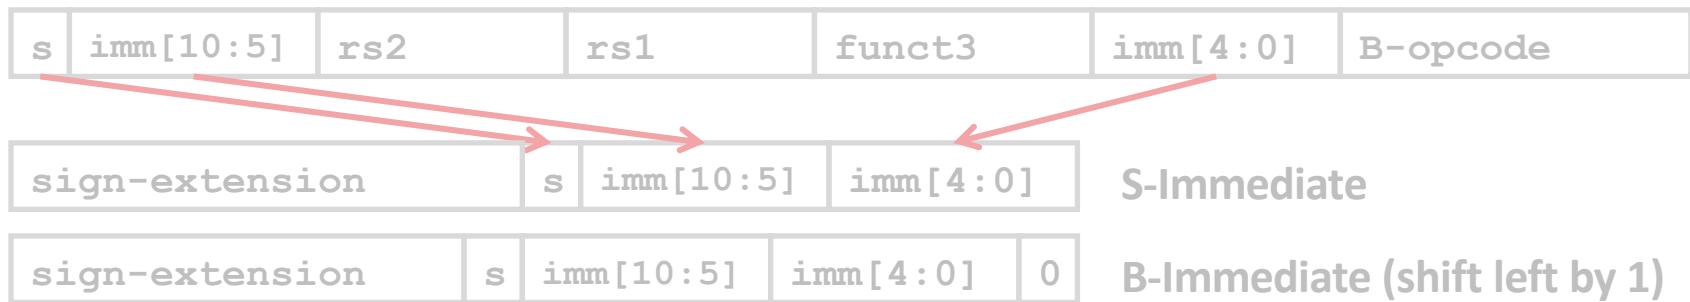
- 12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes
- Standard approach: Treat immediate as in range -2048..+2047, then shift left by 1 bit to multiply by 2 for branches



Each instruction immediate bit can appear in one of two places in output immediate value – so need one 2-way mux per bit

RISC-V Branch Immediates

- 12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes
- RISC-V approach: keep 11 immediate bits in fixed position in output value, and rotate LSB of S-format to be bit 12 of B-format



Only one bit changes position between S and B, so only need a single-bit 2-way mux:

RISC-V Immediate Encoding

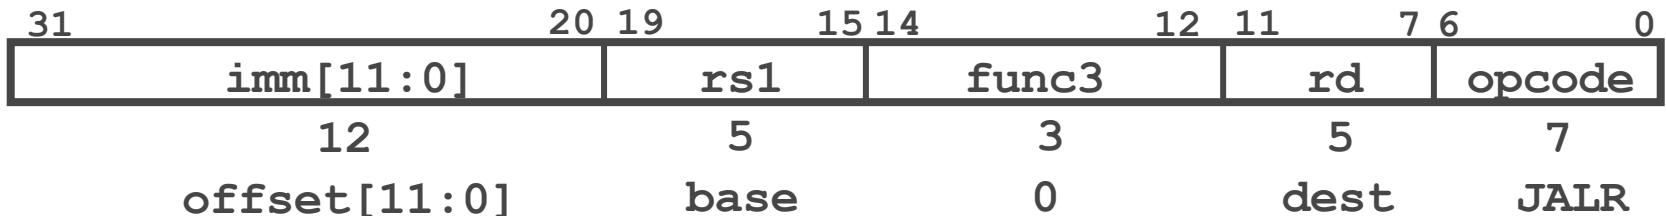
Instruction encodings, $\text{inst}[31:0]$											
31	30	25 24	20 19	15 14	12 11	8	7	6	0		
funct7	rs2	rs1	funct3	rd	opcode						R-type
imm[11:0]		rs1	funct3	rd	opcode						I-type
imm[11:5]	rs2	rs1	funct3	imm[4:0]	opcode						S-type
imm[12 10:5]	rs2	rs1	funct3	imm[4:1 11]	opcode						B-type

32-bit immediates produced, $\text{imm}[31:0]$											
31	25 24	12	11	10	5	4	1	0			
- $\text{inst}[31]$ -		$\text{inst}[30:25]$	$\text{inst}[24:21]$	$\text{inst}[20]$							I-imm.
- $\text{inst}[31]$ -		$\text{inst}[30:25]$	$\text{inst}[11:8]$	$\text{inst}[7]$							S-imm.
- $\text{inst}[31]$ -	$\text{inst}[7]$	$\text{inst}[30:25]$	$\text{inst}[11:8]$					0			B-imm.

Upper bits sign-extended from $\text{inst}[31]$ always

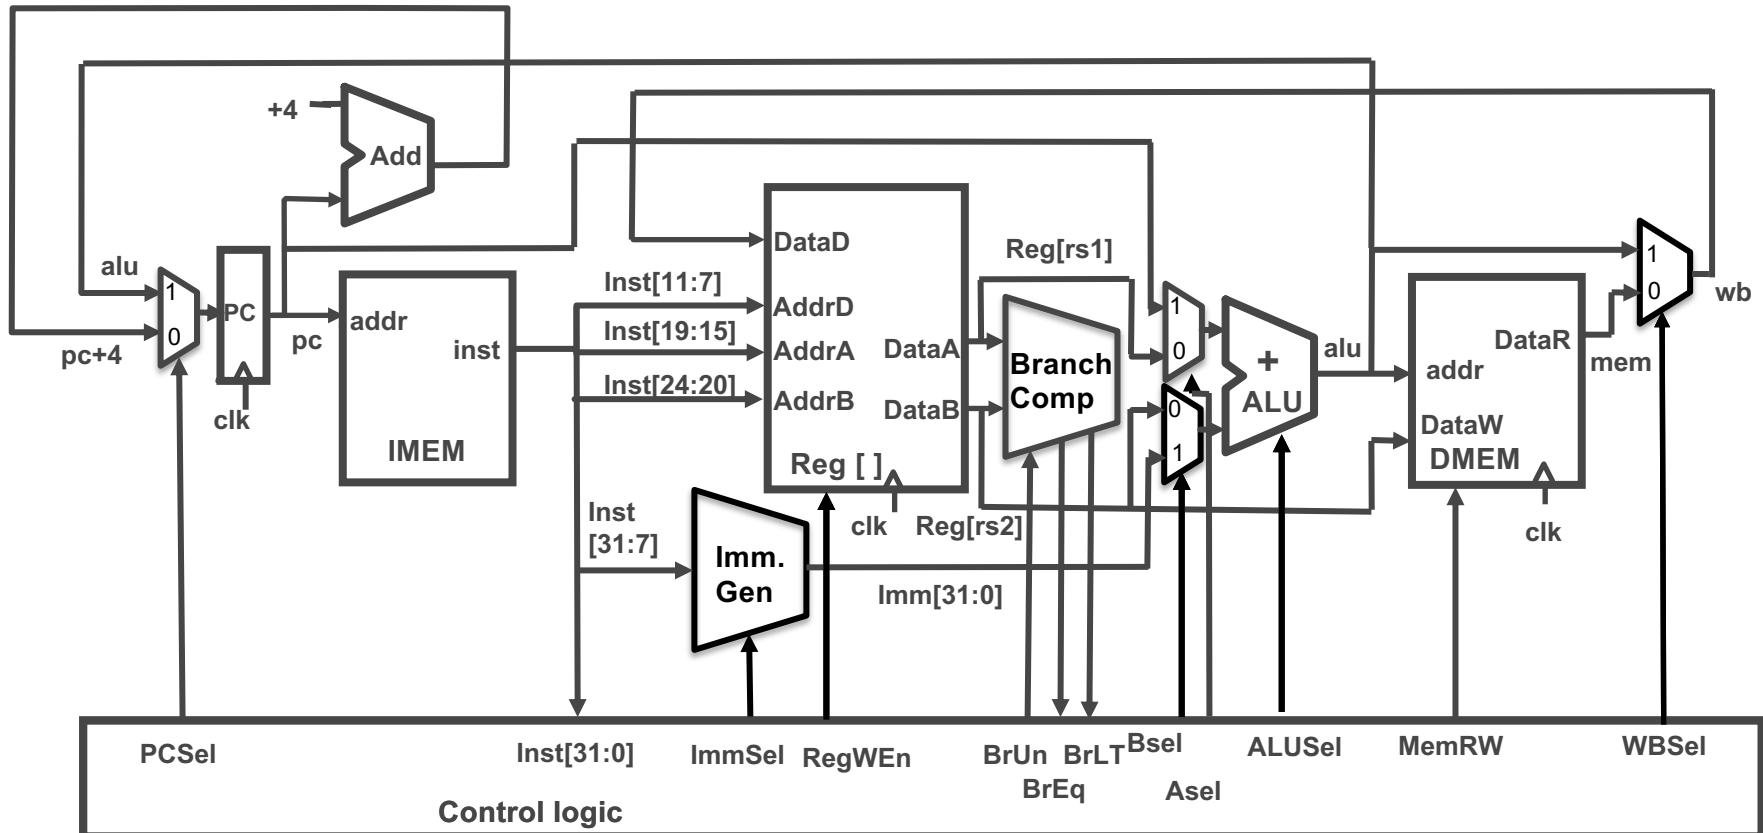
Only bit 7 of instruction changes role in immediate between S and B

Let's Add JALR (I-Format)

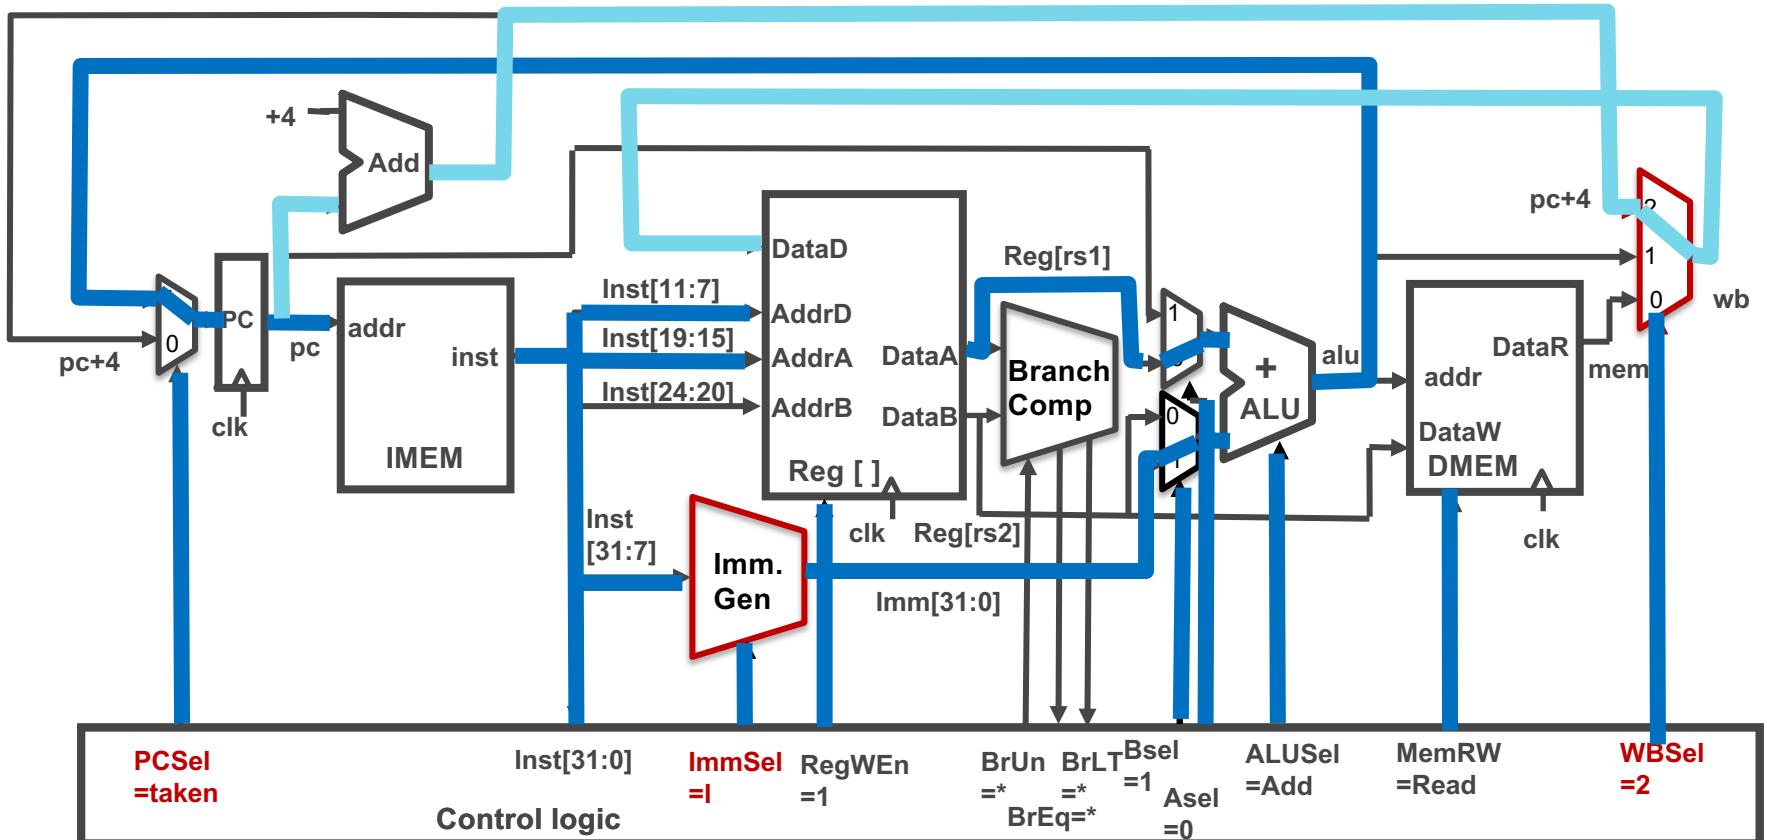


- JALR rd, rs, immediate
- Two changes to the state
 - Writes PC+4 to rd (return address)
 - Sets PC = rs + immediate
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes
 - LSB is ignored

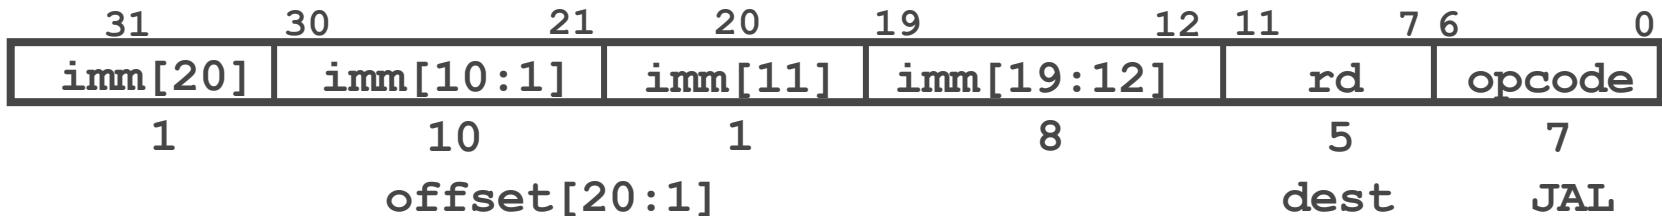
Datapath So Far, with Branches



Adding JALR

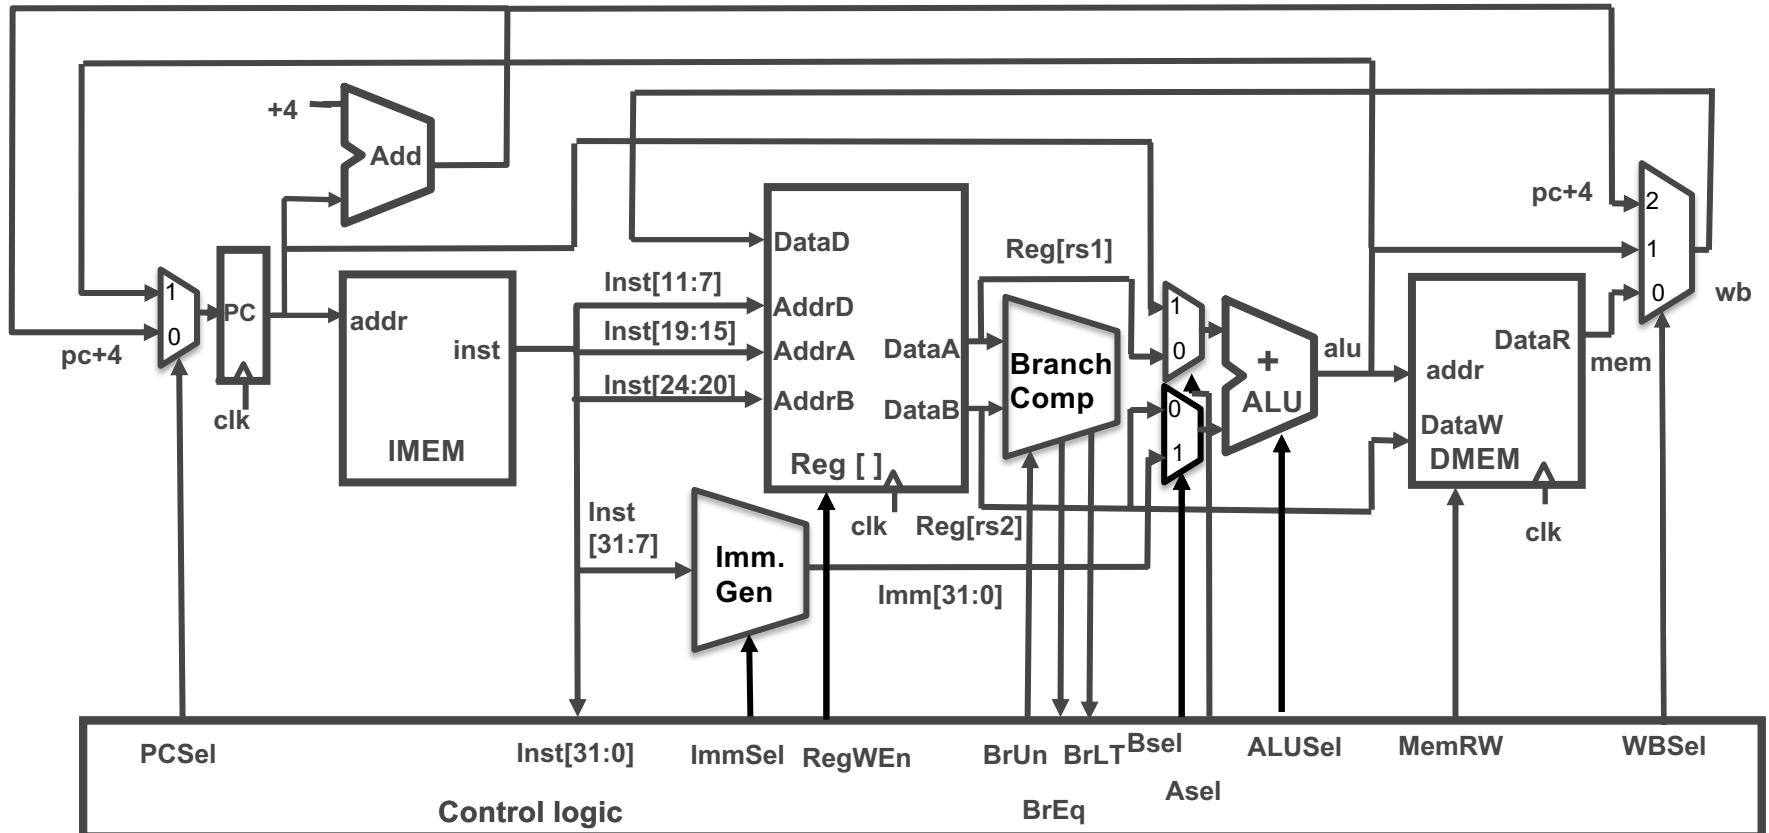


Adding JAL

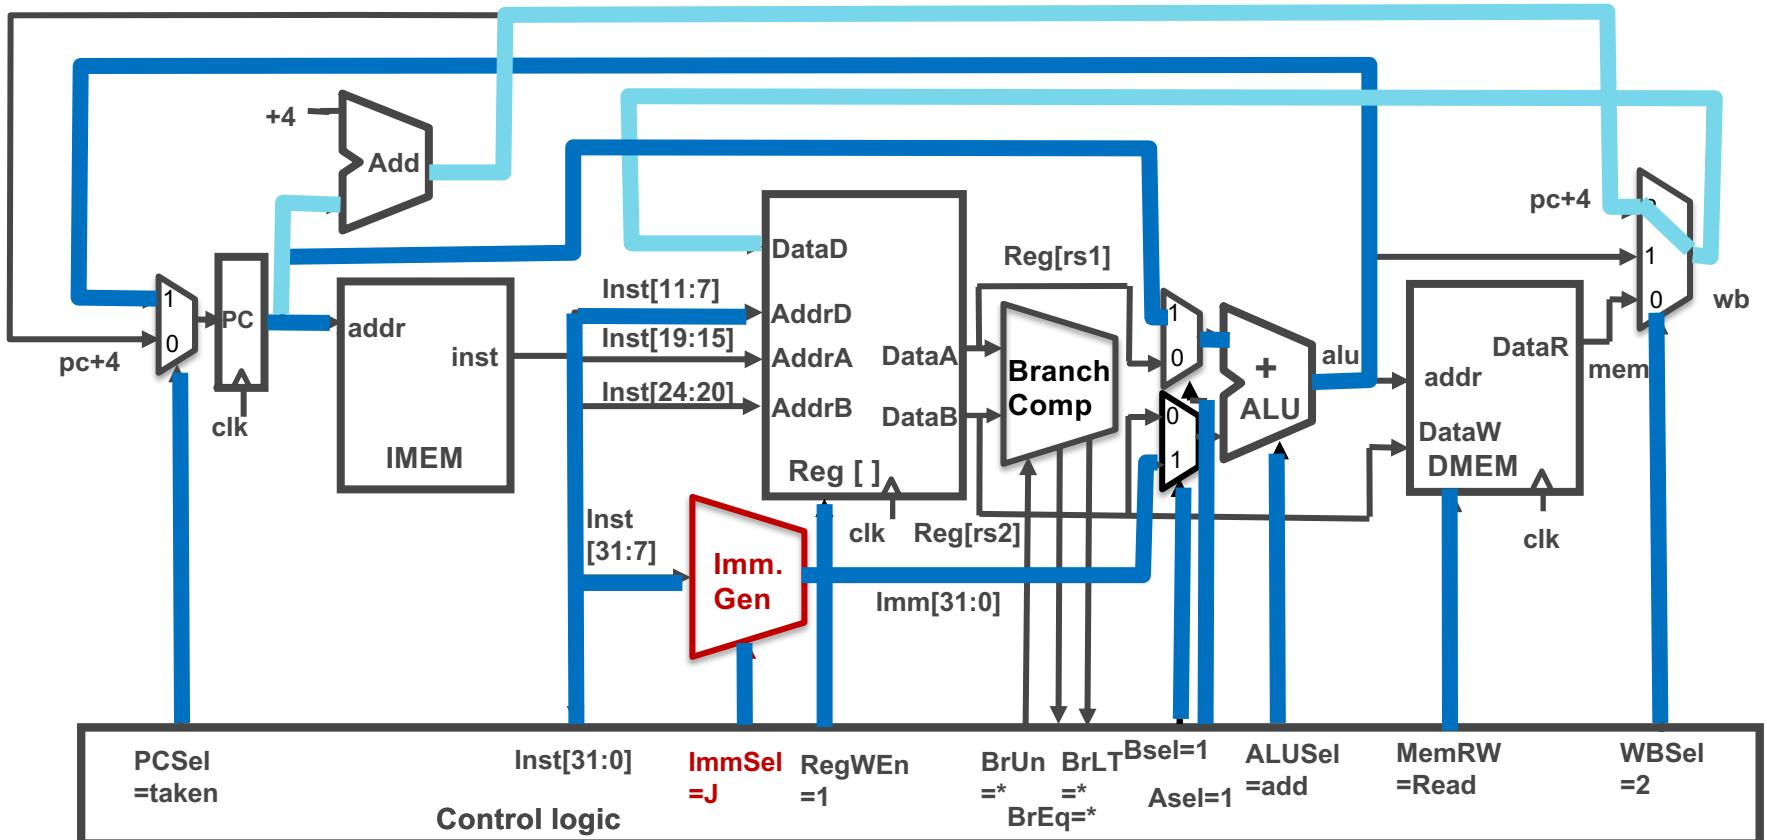


- JAL saves PC+4 in register rd (the return address)
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

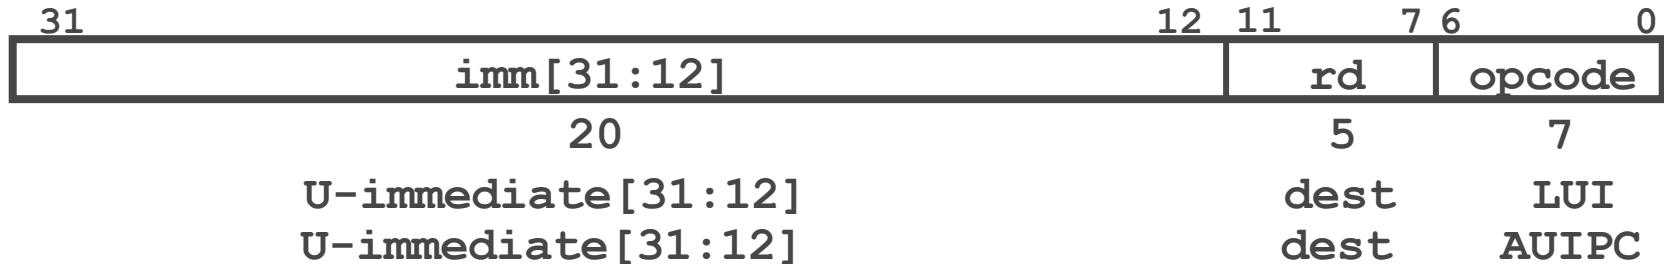
Datapath with JALR



Adding JAL

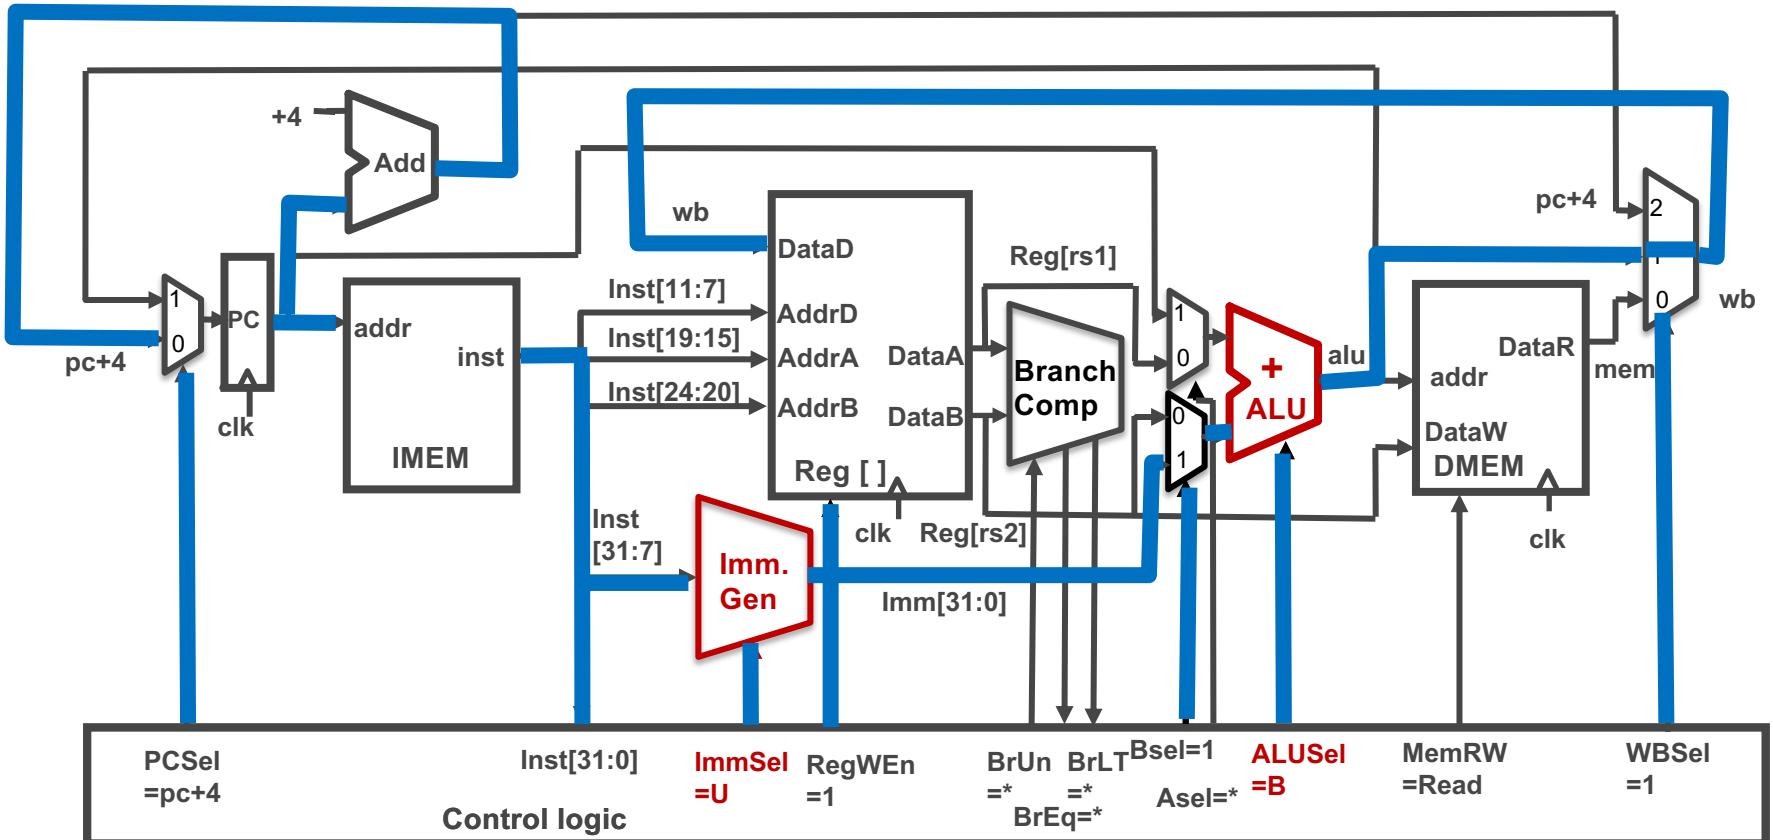


U-Format for “Upper Immediate” Instructions

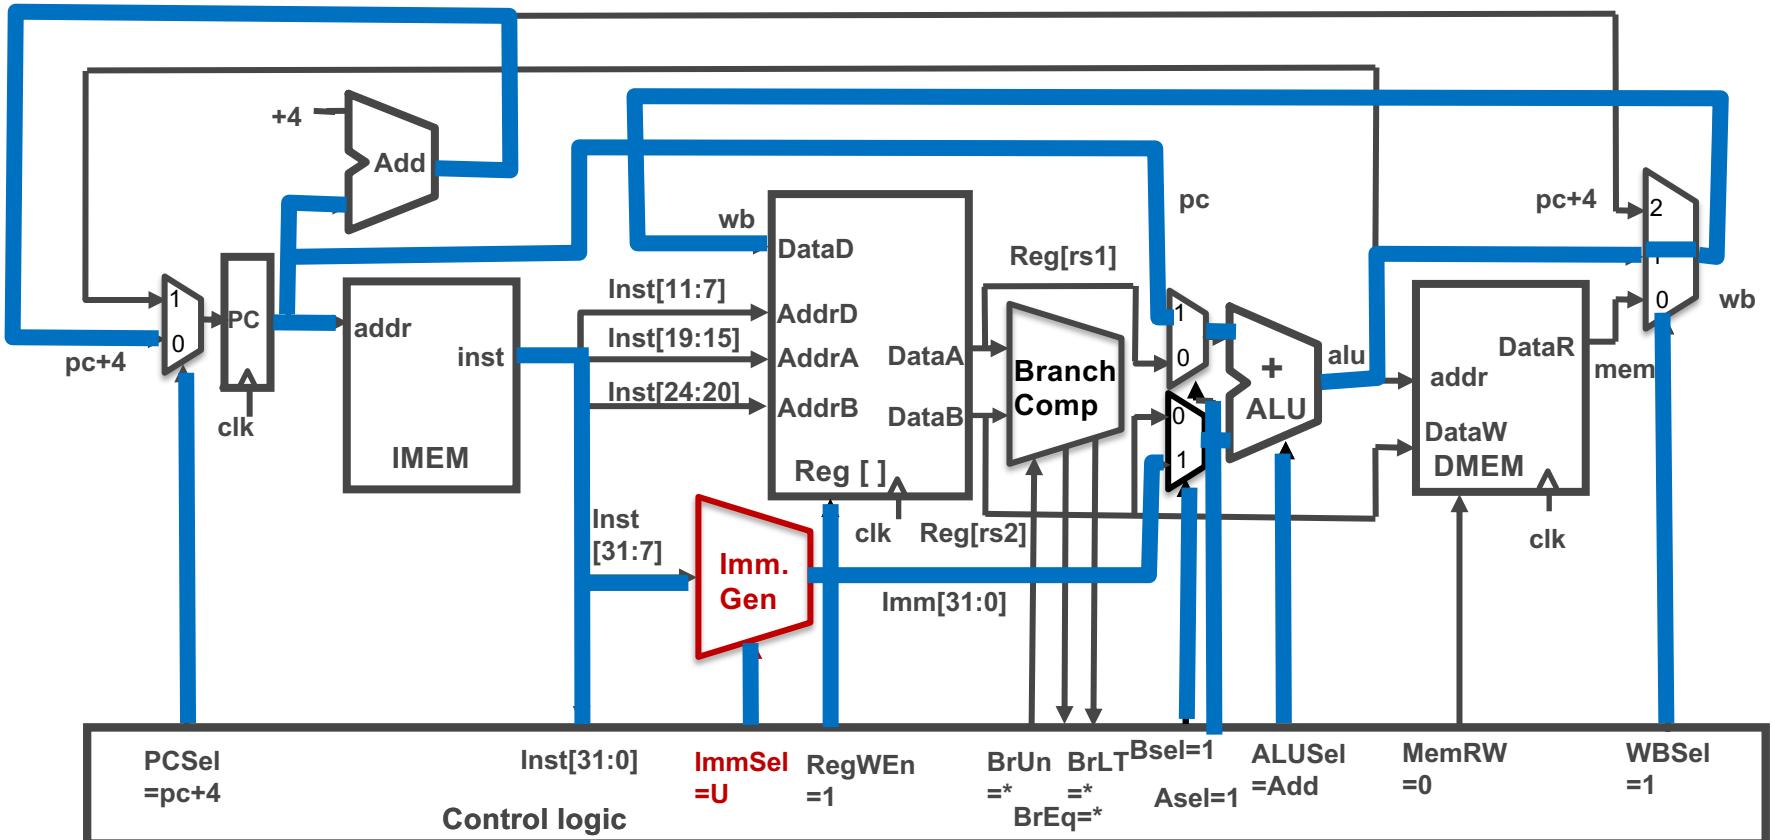


- Has 20-bit immediate in upper 20 bits of 32-bit instruction word
- One destination register, rd
- Used for two instructions
 - LUI – Load Upper Immediate
 - AUIPC – Add Upper Immediate to PC

Implementing LUI



Implementing AUI PC



Recap: Complete RV32I ISA

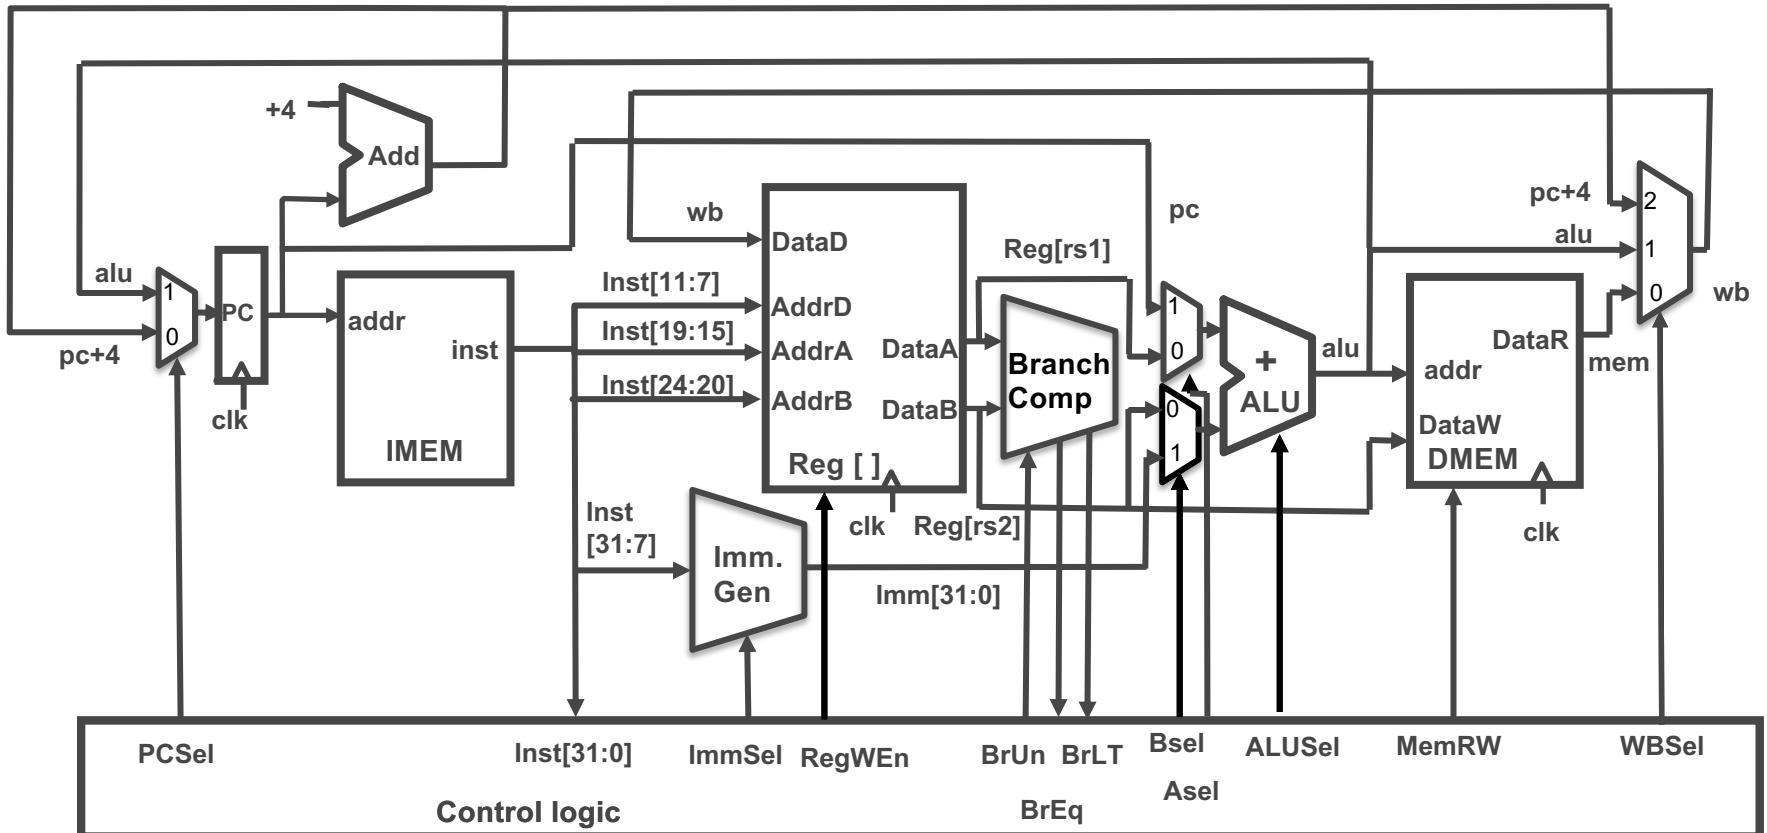
imm[31:12]			rd	0110111
imm[31:12]			rd	0010111
imm[20:10:1 11 19:12]			rd	1101111
imm[11:0]	rs1	000	rd	1100111
imm[12:10:5]	rs2	rs1	000	imm[4:1 11]
imm[12:10:5]	rs2	rs1	001	imm[4:1 11]
imm[12:10:5]	rs2	rs1	100	imm[4:1 11]
imm[12:10:5]	rs2	rs1	101	imm[4:1 11]
imm[12:10:5]	rs2	rs1	110	imm[4:1 11]
imm[12:10:5]	rs2	rs1	111	imm[4:1 11]
imm[11:0]		rs1	000	rd
imm[11:0]		rs1	001	rd
imm[11:0]		rs1	010	rd
imm[11:0]		rs1	100	rd
imm[11:0]		rs1	101	rd
imm[11:5]	rs2	rs1	000	imm[4:0]
imm[11:5]	rs2	rs1	001	imm[4:0]
imm[11:5]	rs2	rs1	010	imm[4:0]
imm[11:0]		rs1	000	rd
imm[11:0]		rs1	010	rd
imm[11:0]		rs1	011	rd
imm[11:0]		rs1	100	rd
imm[11:0]		rs1	110	rd
imm[11:0]		rs1	111	rd

LUI	0000000	shamt	rs1	001	rd	0010011
AUIPC	0000000	shamt	rs1	101	rd	0010011
JAL	0100000	shamt	rs1	101	rd	0010011
JALR	0000000	rs2	rs1	000	rd	0110011
BEQ	0100000	rs2	rs1	000	rd	0110011
BNE	0000000	rs2	rs1	001	rd	0110011
BLT	0000000	rs2	rs1	010	rd	0110011
BGE	0000000	rs2	rs1	011	rd	0110011
BLTU	0000000	rs2	rs1	100	rd	0110011
BGEU	0000000	rs2	rs1	101	rd	0110011
LB	0100000	rs2	rs1	101	rd	0110011
LH	0000000	rs2	rs1	110	rd	0110011
LW	0000000	rs2	rs1	111	rd	0110011
LBU	0000	pred	succ	00000	000	0001111
LHU	0000	0000	0000	00000	001	0001111
SB	00000000000000			00000	000	00000
SH	000000000001			00000	000	00000
SW	csr			rs1	001	rd
ADDI	csr			rs1	010	rd
SLTI	csr			rs1	011	rd
SLTIU	csr			zimm	101	rd
XORI	csr			zimm	110	rd
ORI	csr			zimm	111	rd
ANDI	csr					1110011

Not in CA

- RV32I has 47 instructions
- 37 instructions are enough to run any C program

Complete Single-Cycle RV32I Datapath!



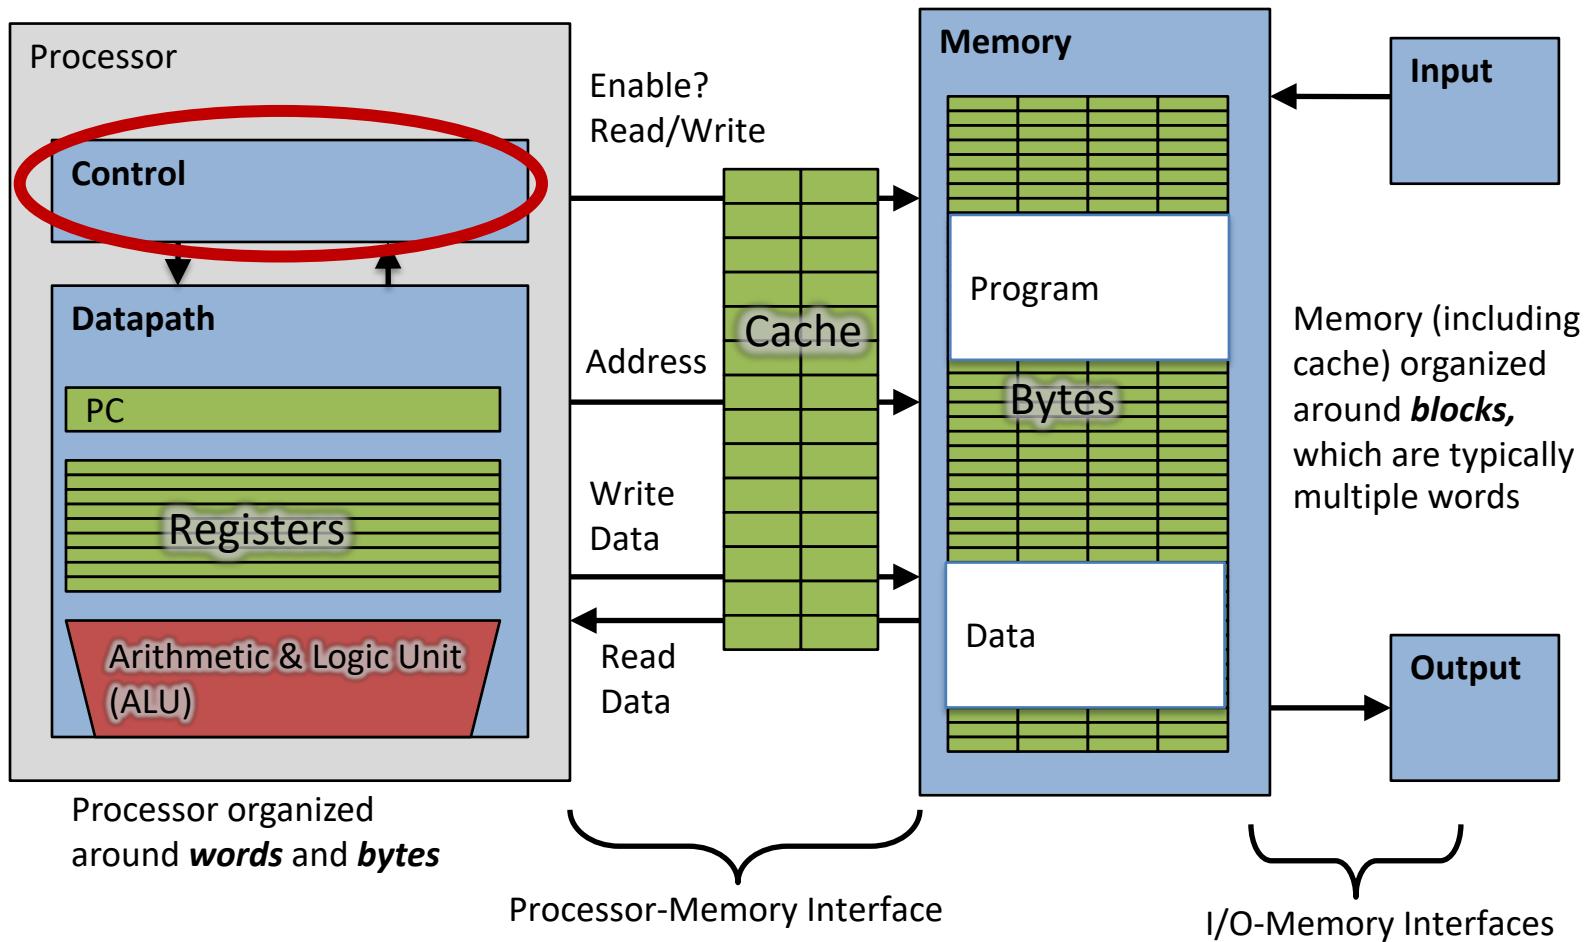
Register Transfer Level (RTL)

- RTL describes instructions in figures or text
- Can use C (or Verilog) to describe RTL

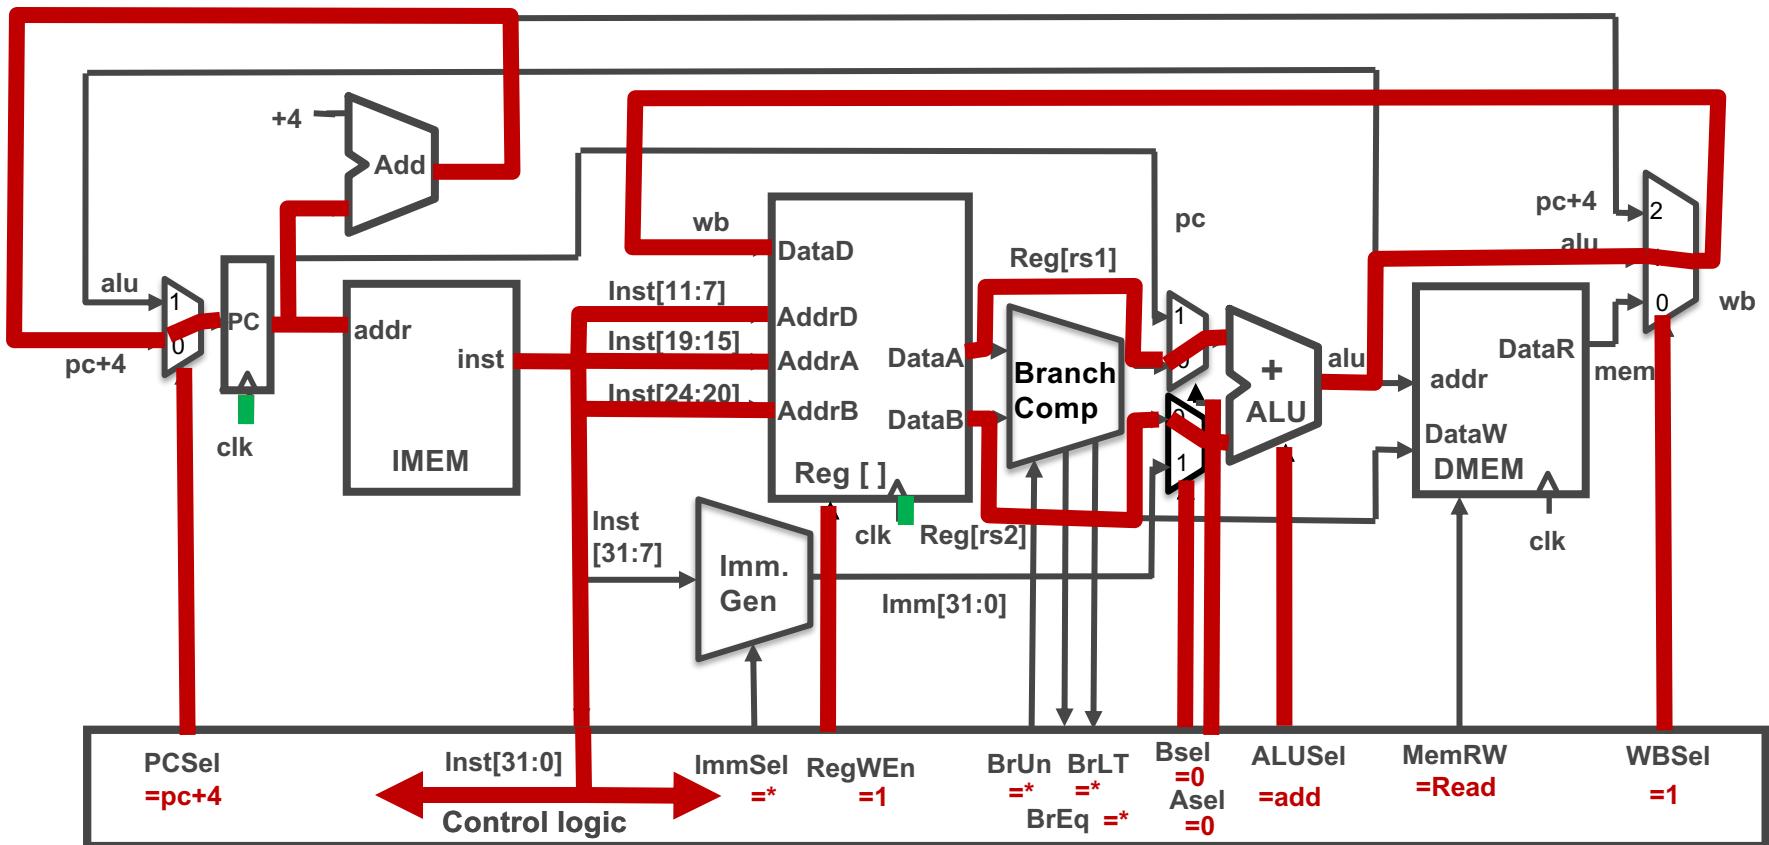
Inst Register Transfers

```
add    R[rd] ← R[rs1] + R[rs2]; PC ← PC + 4
sub    R[rd] ← R[rs1] - R[rs2]; PC ← PC + 4
ori    R[rd] ← R[rs1] | Imm; PC ← PC + 4
jal    R[rd] ← PC + 4; PC ← R[rs1] + Imm
beq    if ( R[rs1] == R[rs2] )
            PC ← PC + Imm
        else PC ← PC + 4
```

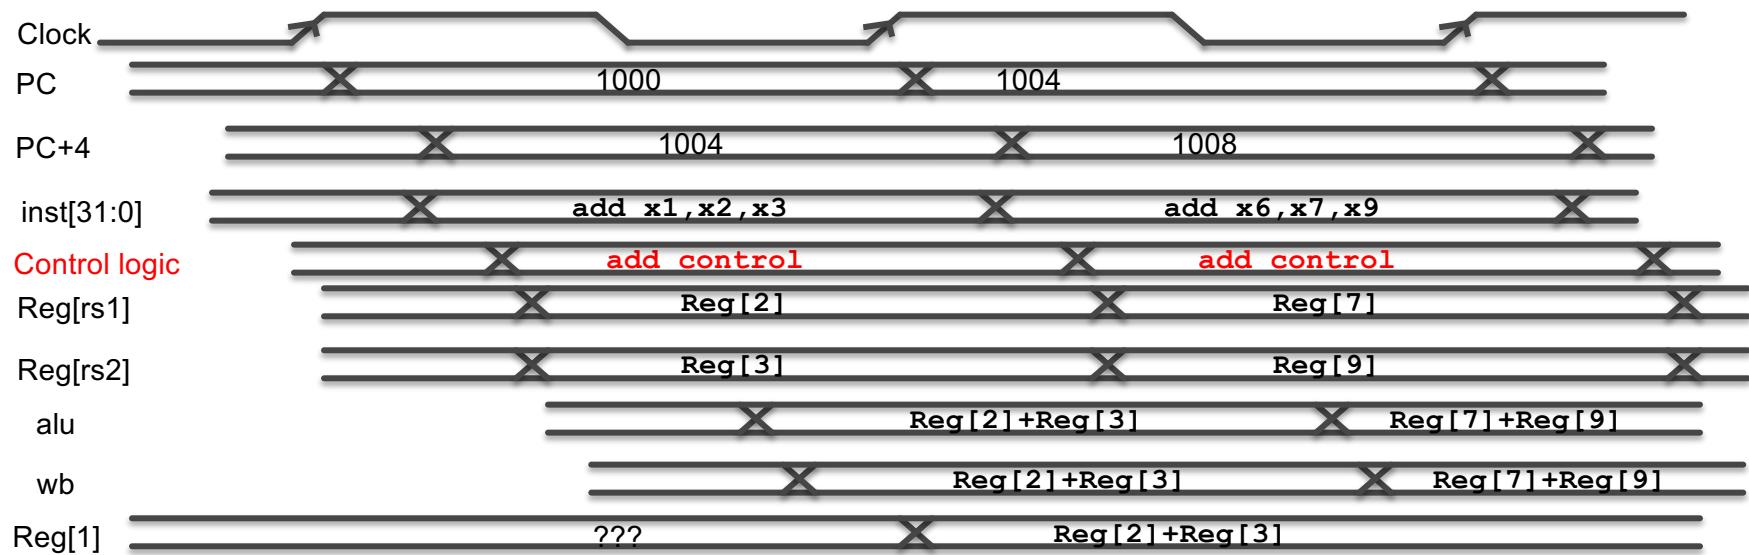
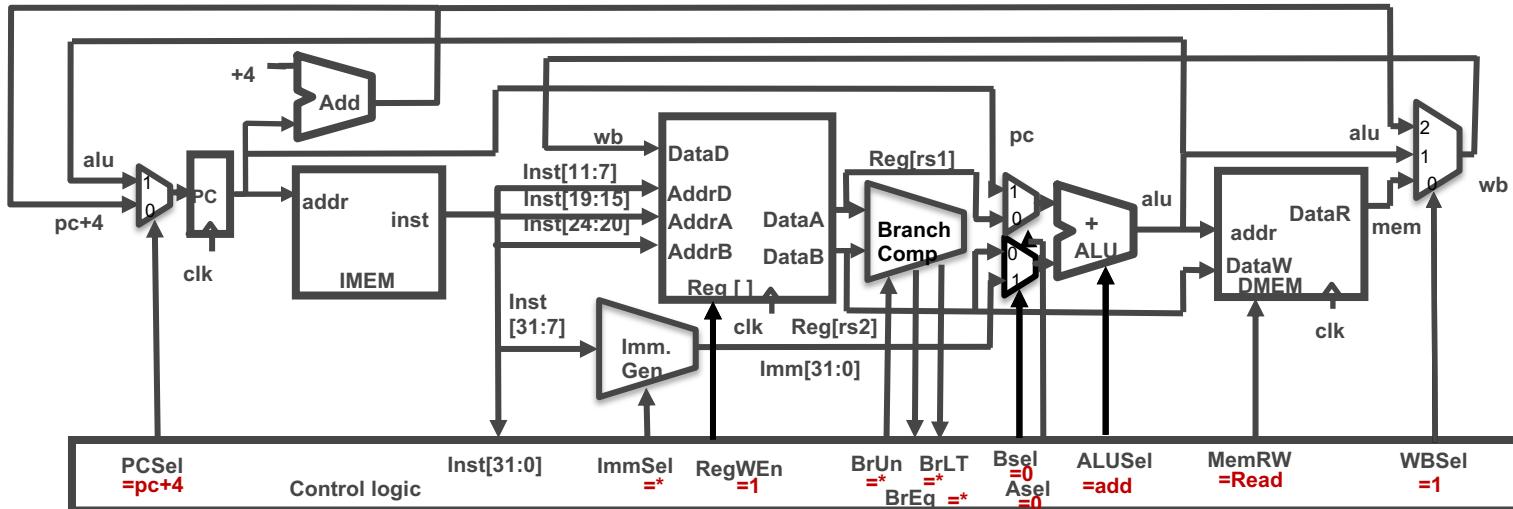
Our Single-Core Computer (without FPU, SIMD)



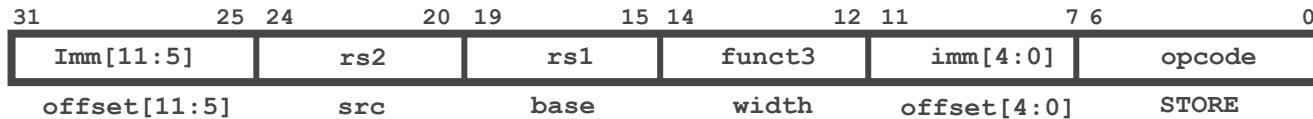
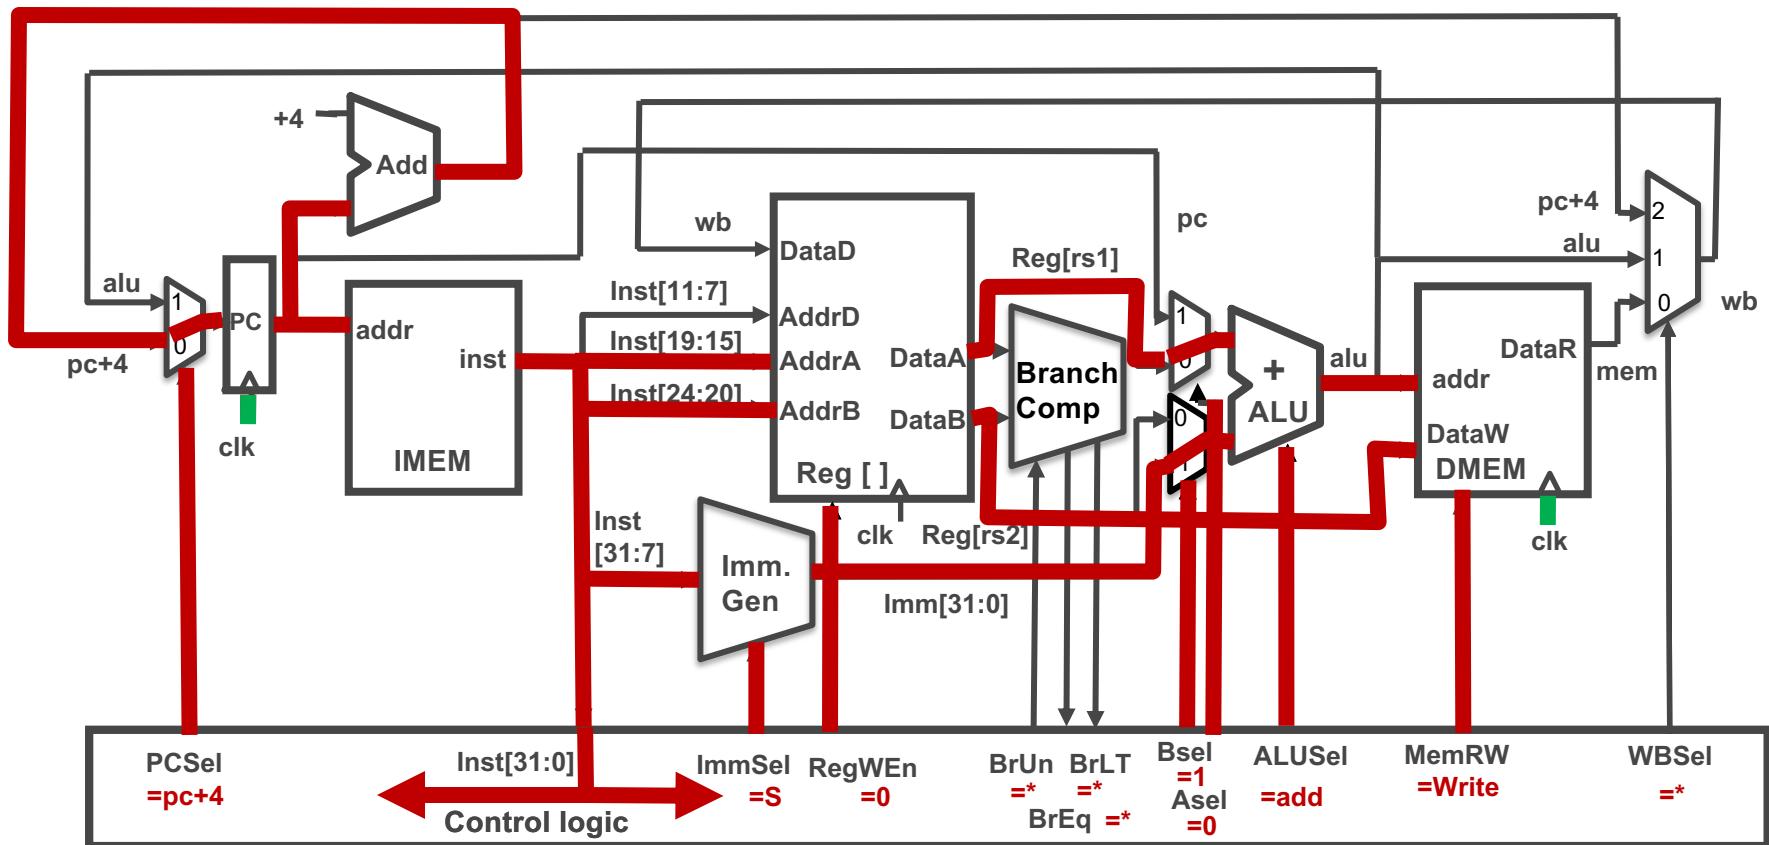
Example: add



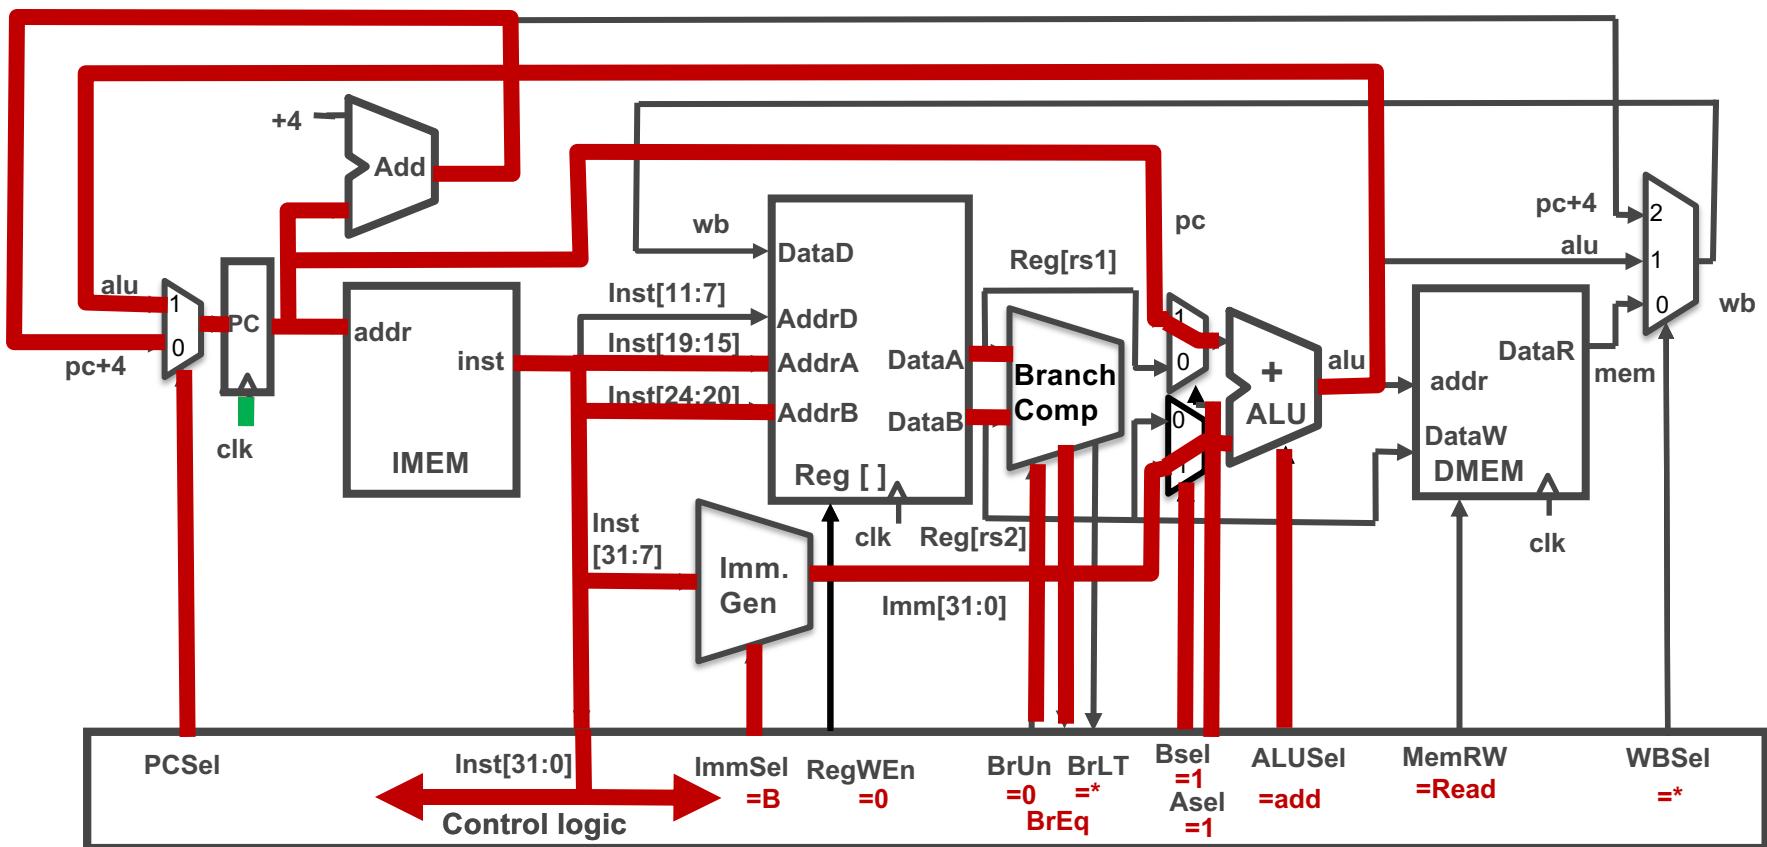
add Execution



Example: SW



Example: beq



Midterm I

- Date: Wednesday, Apr. 03
- Time: 08:15- 9:55 (normal lecture slot)
 - Be punctual – we start 8:15 sharp!
- Venue: Teaching Center 301 + 201
- Closed book:
 - You can bring one A4 page with notes (both sides; in English): Write your Chinese and **Pinyin** name on the top! **Handwritten** by you!
 - You will be provided with the RISC-V "green sheet"
 - No other material allowed!

Midterm I

- Switch cell phones **off!** (not silent mode – off!)
 - Put them in your bags.
- Bags under the table. Nothing except paper, pen, 1 drink, 1 snack on the table!
- No other electronic devices are allowed!
 - No ear plugs, music, smartwatch...
- Anybody touching any electronic device will **FAIL** the course!
- Anybody found cheating (copy your neighbors answers, additional material, ...) will **FAIL** the course!









Old School Machine Structures

- C program starts
- C executable and is loaded into memory by OS (copying bytes)
- OS sets up stack, then tells sets up stack
- Run time first initializes memory and other libraries
- Then will your procedure run and main

New School Machine Structures

- Parallel Requests
- Parallel Threads
- Parallel Data
- Hardware descriptors

6 Great Ideas in Computer Architecture

1. Abstraction
2. Levels of Representation / Interpretation
3. Moore's Law
4. Decoupling through threads
5. Principle of Locality
6. Memory Hierarchy
7. Parallelism
8. Performance Measurement and Improvement
9. Dependability via Redundancy

Two's Complement Representation

- treats 0 as positive
- 32-bit word represents 2^{32} states from -2^{31} to $2^{31}-1$

Components of a Computer

fw. c → cpp → fo.1 → compiler

- cpp replaces comments with a single space
- cpp commands begins with '#'

(standard) Type

- (signed) char / signed char
- (unsigned) int 4
- (signed) short 2 + long × 2
- (unsigned) long 4
- float 4 byte
- double 8 byte

Level of Representation / Interpretation

10. 指令集 1. 指令
11. Q Q 是指与
12. II 逻辑或
13. ?: 条件左
14. = + = Q =
- = - = 1 = to → to
- <= <= 1 = to → to
- >= >= 1 = to → to
15. , to → to
- long long 32 8
- long long 64 8

Valid Pointer Arithmetic

- Add an integer to a pointer
- Subtract 2 pointers (in the same array)
- Compare two pointers ($C, C+1, \dots$)
- Run time first initializes memory and other libraries
- Run time first initializes memory and other libraries
- Then will your procedure run and main

Six Fundamental Steps in Calling a Function

1. Put parameters in a place where function can access them
2. Transfer control to function
3. Acquire local storage resources needed for function
4. Perform desired task of the function
5. Put result value in a place where calling code can access it and restore any registers you used
6. Return control to point of origin, since a function can be called from several parts of a program.

\$fp

Saved argument register
Saved return address
Saved saved registers
Local arrays and structures

sp

Stack
↓
Dynamic data
Saved static data
Temporary
↓
Local arrays and structures

bp

Stack
↓
Dynamic data
Saved static data
Temporary
↓
Local arrays and structures

Code

Stack

- free when function returns
- Last in first out

Format

- used for instructions with immediates, lw and sw and branches (beq and bne)

Common Memory Problems

- Using uninitialized variables
- Using memory that you don't own
- Improper use of free / realloc by memory with the function handle returned by malloc / calloc
- Memory leaks

Opcode rs rt rd shamt funct

opcode = 0

2 format: 5-bit field only represents numbers up to 31

1st instruction has immediate, then it uses at most 2 registers used for lw, sw, beq, bne branches and with immediate

Decoding with larger immediate

- Label Upper lui
- add \$to \$to \$to \$to \$to \$to \$to
- lui \$tut \$tut \$tut \$tut
- ori \$tut \$tut \$tut \$tut
- addu \$to \$to \$to \$to

Architectural Implementation

Large Constant Descriptor

THE PURPOSE OF THIS REFERENCE GUIDE IS TO WALK THROUGH THE PROCESS OF BOOTING THE SIFT WORKSTATION, CREATING A TIMELINE ("SUPER" OR "MICRO") AND REVIEWING IT.



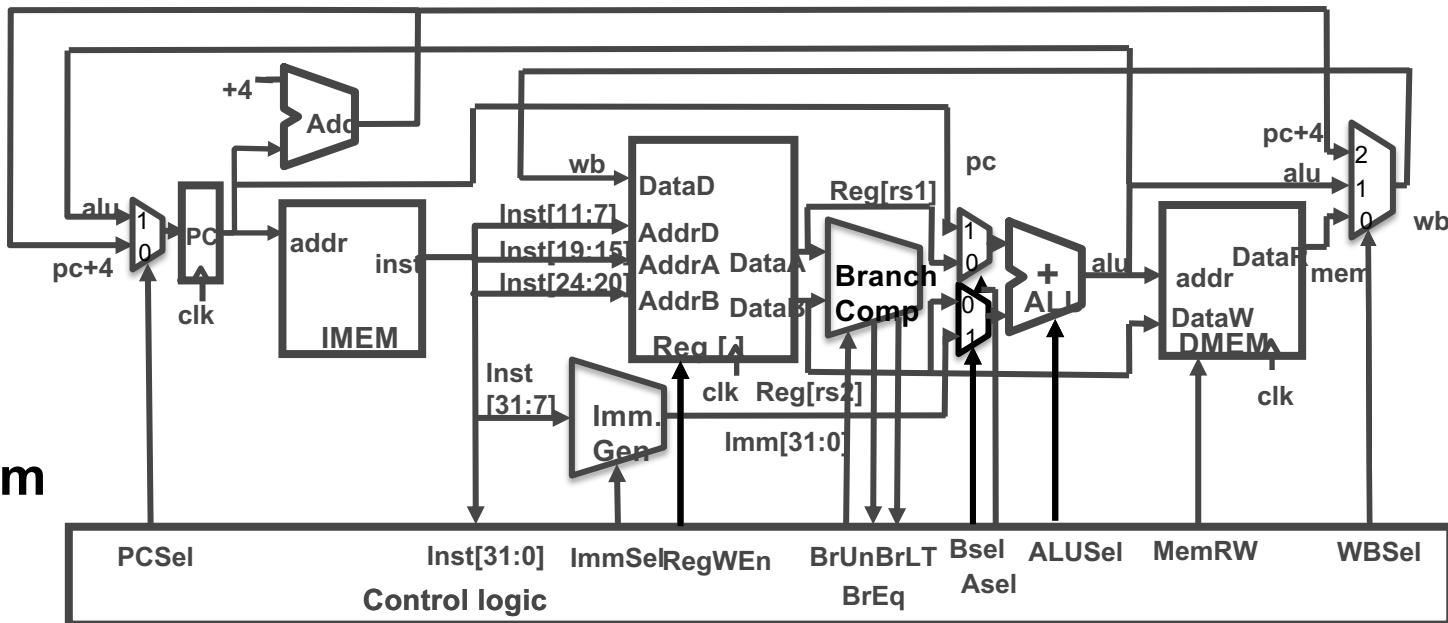
Content

- Main topics
 - Number representation
 - C
 - RISC-V
- Plus general “Computer Architecture” knowledge
- Everything till lecture 8 CALL – including lecture 8

Peer Instruction: Critical Path

Critical path for addi

$$R[rd] = R[rs1] + imm$$

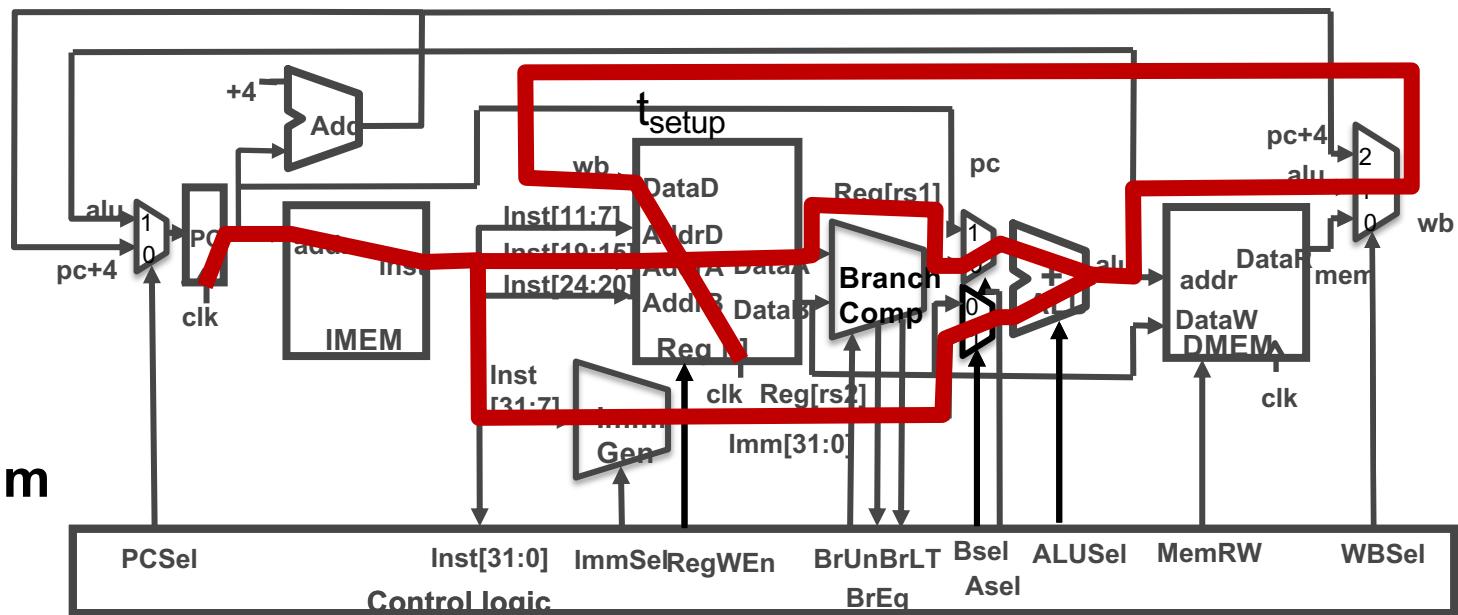


- A. $t_{clk-q} + t_{Add} + t_{IMEM} + t_{Reg} + t_{BComp} + t_{ALU} + t_{DMEM} + t_{mux} + t_{Setup}$
- B. $t_{clk-q} + t_{IMEM} + \max\{t_{Reg}, t_{Imm}\} + t_{ALU} + 2t_{mux} + t_{Setup}$
- C. $t_{clk-q} + t_{IMEM} + \max\{t_{Reg}, t_{Imm}\} + t_{ALU} + 3t_{mux} + t_{DMEM} + t_{Setup}$
- D. None of the above

Peer Instruction: Critical Path

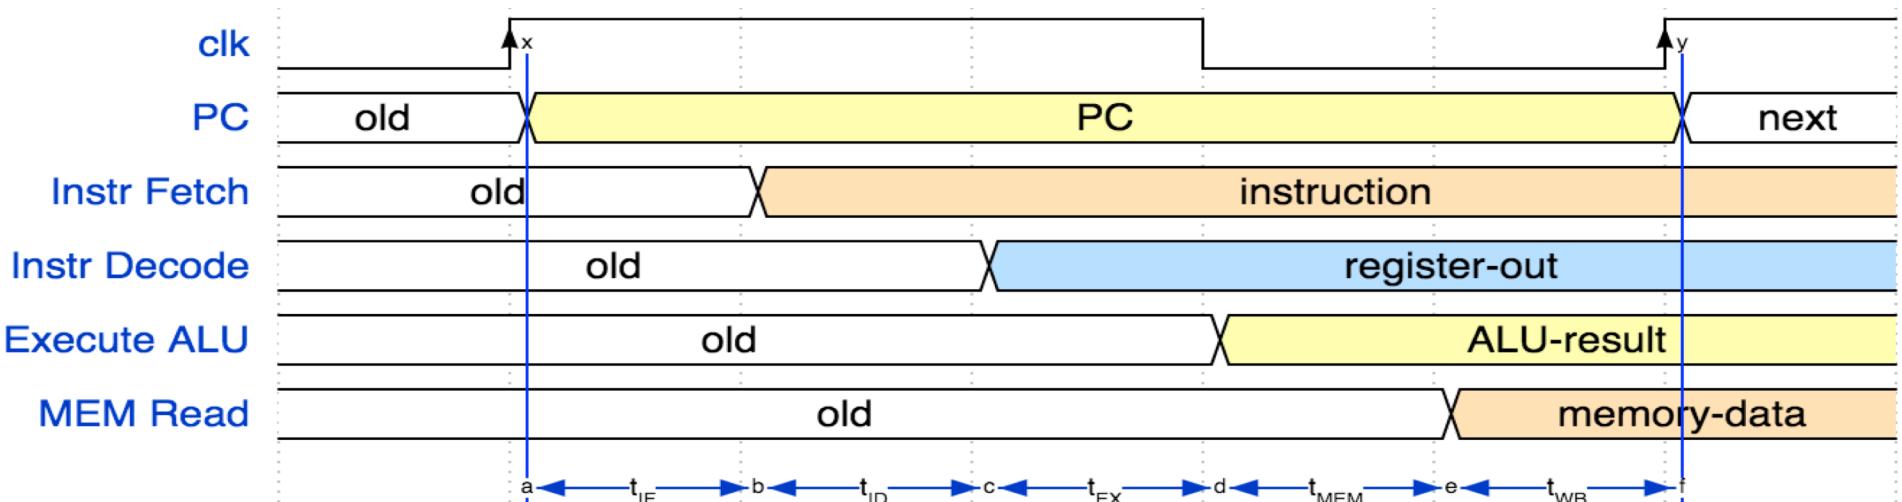
Critical path for addi

$$R[rd] = R[rs1] + imm$$



- A. $t_{clk-q} + t_{Add} + t_{IMEM} + t_{Reg} + t_{BComp} + t_{ALU} + t_{DMEM} + t_{mux} + t_{Setup}$
- B. $t_{clk-q} + t_{IMEM} + \max\{t_{Reg}, t_{Imm}\} + t_{ALU} + 2*t_{mux} + t_{Setup}$
- C. $t_{clk-q} + t_{IMEM} + \max\{t_{Reg}, t_{Imm}\} + t_{ALU} + 3*t_{mux} + t_{DMEM} + t_{Setup}$
- D. None of the above

Instruction Timing



IF	ID	EX	MEM	WB	Total
I-MEM	Reg Read	ALU	D-MEM	Reg W	
200 ps	100 ps	200 ps	200 ps	100 ps	800 ps

Instruction Timing

Instr	IF = 200ps	ID = 100ps	ALU = 200ps	MEM=200ps	WB = 100ps	Total
add	X	X	X		X	600ps
beq	X	X	X			500ps
jal	X	X	X		X	600ps
lw	X	X	X	X	X	800ps
sw	X	X	X	X		700ps

- Maximum clock frequency
 - $f_{\max} = 1/800\text{ps} = 1.25 \text{ GHz}$
- Most blocks idle most of the time
 - E.g. $f_{\max, \text{ALU}} = 1/200\text{ps} = 5 \text{ GHz!}$

Control Logic Truth Table

Inst[31:0]	BrEq	BrLT	PCSel	ImmSel	BrUn	ASel	BSel	ALUSel	MemRW	RegWEn	WBSel
add	*	*	+4	*	*	Reg	Reg	Add	Read	1	ALU
sub	*	*	+4	*	*	Reg	Reg	Sub	Read	1	ALU
(R-R Op)	*	*	+4	*	*	Reg	Reg	(Op)	Read	1	ALU
addi	*	*	+4	I	*	Reg	Imm	Add	Read	1	ALU
lw	*	*	+4	I	*	Reg	Imm	Add	Read	1	Mem
sw	*	*	+4	S	*	Reg	Imm	Add	Write	0	*
beq	0	*	+4	B	*	PC	Imm	Add	Read	0	*
beq	1	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	0	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	1	*	+4	B	*	PC	Imm	Add	Read	0	*
blt	*	1	ALU	B	0	PC	Imm	Add	Read	0	*
bltu	*	1	ALU	B	1	PC	Imm	Add	Read	0	*
jalr	*	*	ALU	I	*	Reg	Imm	Add	Read	1	PC+4
jal	*	*	ALU	J	*	PC	Imm	Add	Read	1	PC+4
auiopc	*	*	+4	U	*	PC	Imm	Add	Read	1	ALU ⁴²

Control Realization Options

- ROM
 - “Read-Only Memory”
 - Regular structure
 - Can be easily reprogrammed
 - fix errors
 - add instructions
 - Popular when designing control logic manually
- Combinatorial Logic
 - Today, chip designers use logic synthesis tools to convert truth tables to networks of gates

RV32I, a nine-bit ISA!

imm[31:12]			rd	0110111
imm[31:12]			rd	0010111
imm[20:10:1 11 19:12]			rd	1101111
imm[11:0]	rs1	000	rd	1100111
imm[12:10:5]	rs2	rs1	000	imm[4:1 11]
imm[12:10:5]	rs2	rs1	001	imm[4:1 11]
imm[12:10:5]	rs2	rs1	100	imm[4:1 11]
imm[12:10:5]	rs2	rs1	101	imm[4:1 11]
imm[12:10:5]	rs2	rs1	110	imm[4:1 11]
imm[12:10:5]	rs2	rs1	111	imm[4:1 11]
imm[11:0]		rs1	000	rd
imm[11:0]		rs1	001	rd
imm[11:0]		rs1	010	rd
imm[11:0]		rs1	100	rd
imm[11:0]		rs1	101	rd
imm[11:5]	rs2	rs1	000	imm[4:0]
imm[11:5]	rs2	rs1	001	imm[4:0]
imm[11:5]	rs2	rs1	010	imm[4:0]
imm[11:0]		rs1	000	rd
imm[11:0]		rs1	010	rd
imm[11:0]		rs1	011	rd
imm[11:0]		rs1	100	rd
imm[11:0]		rs1	110	rd
imm[11:0]		rs1	111	rd

inst[30]			inst[14:12] inst[6:2]		
0000000	shamt	rs1	001	rd	0010011
0000000	shamt	rs1	101	rd	0010011
0100000	shamt	rs1	101	rd	0010011
0000000	rs2	rs1	000	rd	0110011
0100000	rs2	rs1	000	rd	0110011
0000000	rs2	rs1	001	rd	0110011
0000000	rs2	rs1	010	rd	0110011
0000000	rs2	rs1	011	rd	0110011
0000000	rs2	rs1	100	rd	0110011
0000000	rs2	rs1	101	rd	0110011
0100000	rs2	rs1	101	rd	0110011
0000000	rs2	rs1	110	rd	0110011
0000000	rs2	rs1	111	rd	0110011
0000000	pred	succ	00000	000	0001111
00000	0000	0000	00000	001	0001111
0000000000000000			00000	000	000000011
0000000000000001			00000	000	000000011
csr		rs1	001	rd	1110011
csr		rs1	010	rd	1110011
csr		rs1	011	rd	1110011
zimm		zimm	101	rd	1110011
zimm		zimm	110	rd	1110011
zimm		zimm	111	rd	1110011

Not in CS61C

Instruction type encoded using only 9 bits
inst[30],inst[14:12], inst[6:2]

Combinational Logic Control

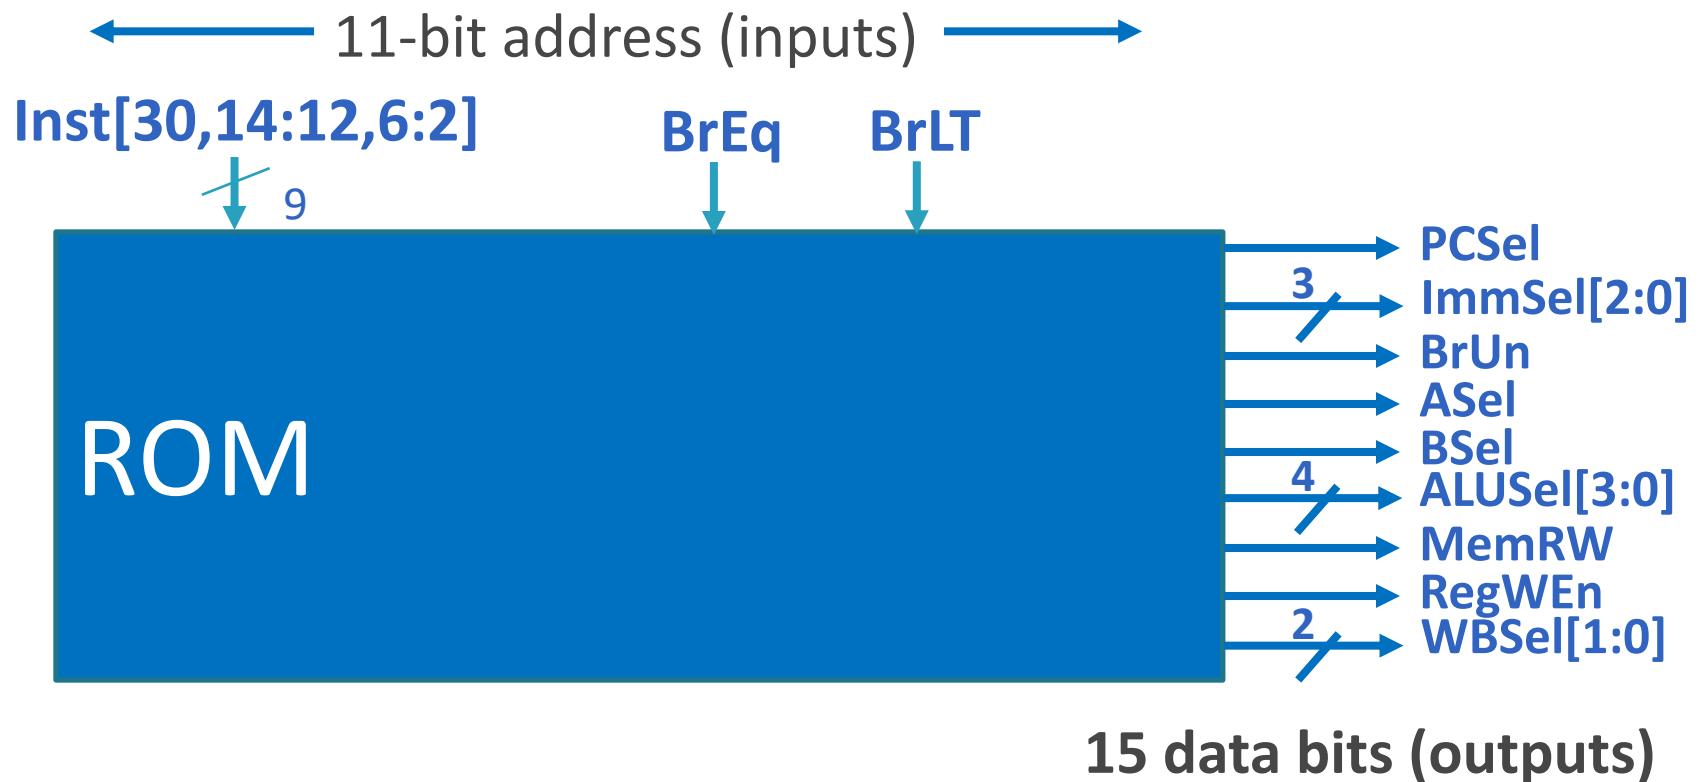
- Simplest example: BrUn

			inst[14:12]	inst[6:2] = Branch		
imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	11000 11	BEQ
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	11000 11	BNE
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	11000 11	BLT
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	11000 11	BGE
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	11000 11	BLTU
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	11000 11	BGEU

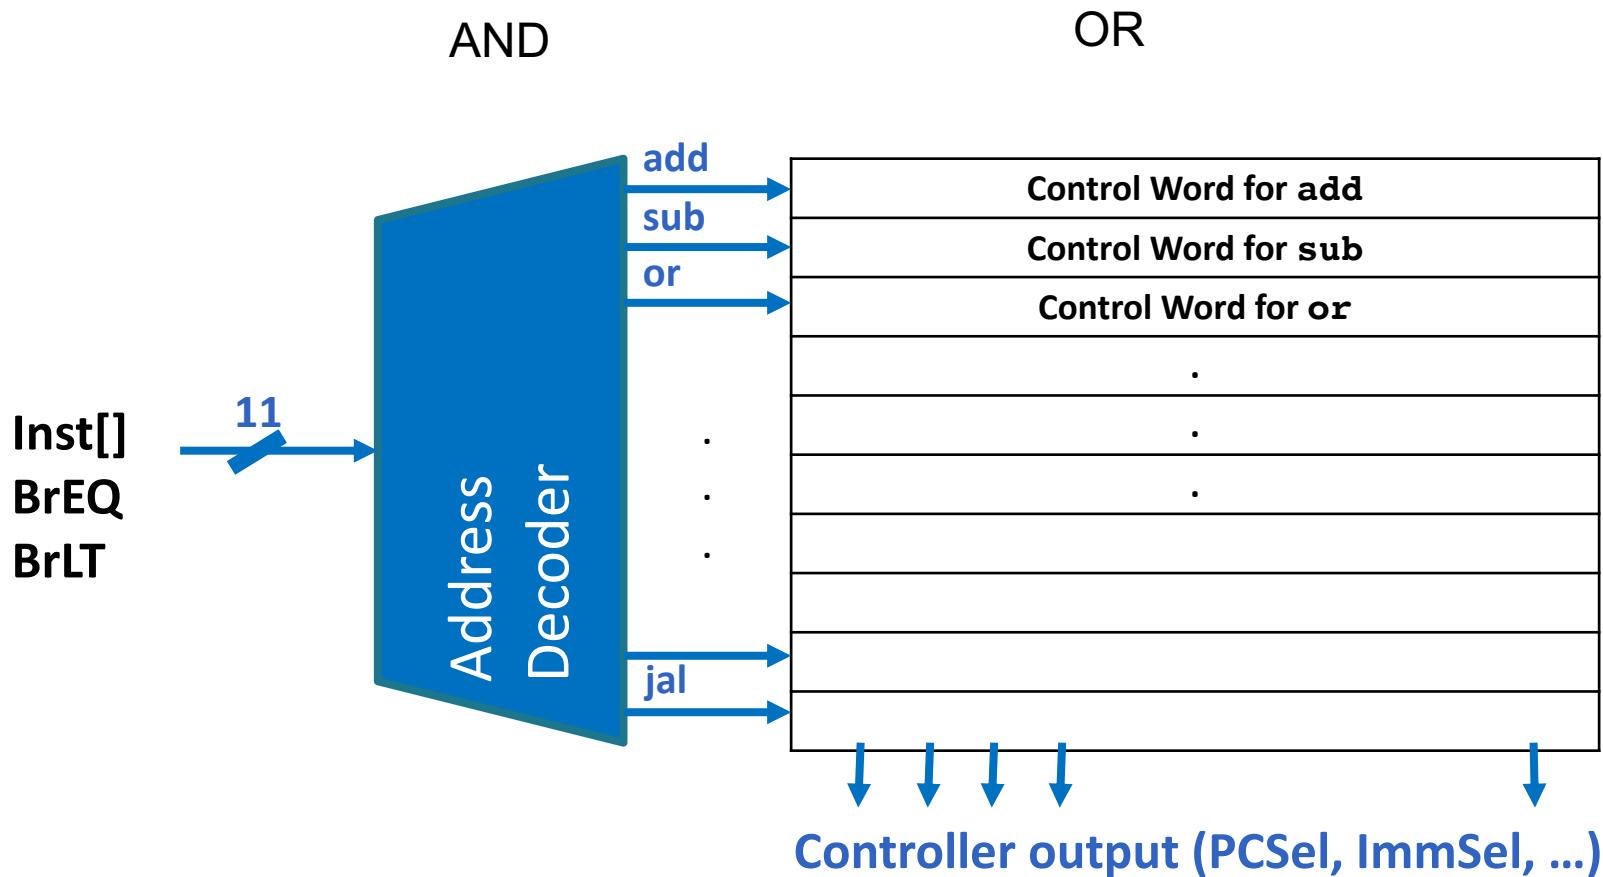
- How to decode whether BrUn is 1?

- BrUn = Inst [13] • Branch

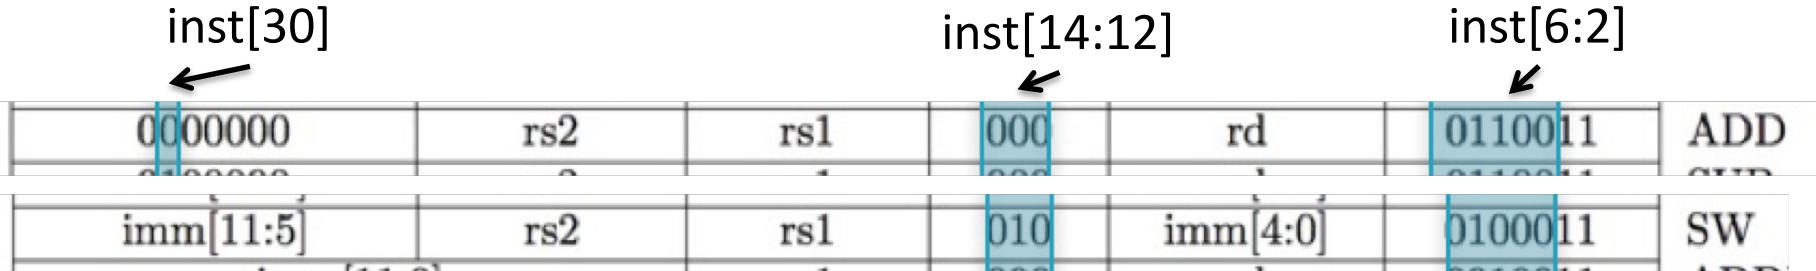
ROM-based Control



ROM Controller Implementation



Controller: AND logic (add, sw)



- $\text{add} = \overline{\text{inst}[2]} \cdot \overline{\text{inst}[3]} \cdot \overline{\text{inst}[4]} \cdot \overline{\text{inst}[5]} \cdot \overline{\text{inst}[6]} \cdot \overline{\text{inst}[12]} \cdot \overline{\text{inst}[13]} \cdot \overline{\text{inst}[14]} \cdot \overline{\text{inst}[30]}$
- $\text{sw} = \overline{\text{inst}[2]} \cdot \overline{\text{inst}[3]} \cdot \overline{\text{inst}[4]} \cdot \overline{\text{inst}[5]} \cdot \overline{\text{inst}[6]} \cdot \overline{\text{inst}[12]} \cdot \overline{\text{inst}[13]} \cdot \overline{\text{inst}[14]}$

AND Controller Logic

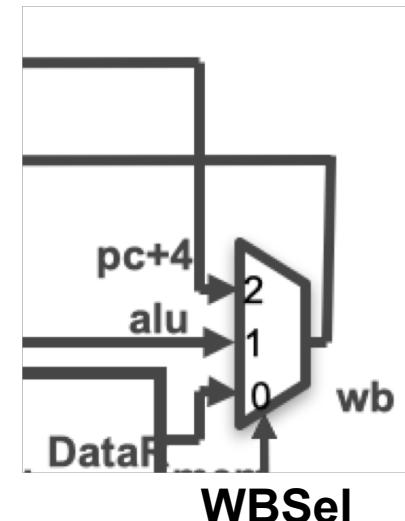
	2	3	4	5	6	12	13	14	30	unused
add	xor(i[2], 1) • xor(i[3], 1) • xor(i[4], 0) • xor(i[5], 0) • xor(i[6], 1) • xor(i[12], 1) • xor(i[13], 1) • xor(i[14], 1) • (xor(i[30], 1) + 0)									
sw	xor(i[2], 1) • xor(i[3], 1) • xor(i[4], 1) • xor(i[5], 0) • xor(i[6], 1) • xor(i[12], 1) • xor(i[13], 1) • xor(i[14], 1) • (xor(i[30], 1) + 1)									
...										

- $\text{bne_t} = \text{bne} \cdot \overline{\text{BrEQ}}$
- $\text{bne_f} = \text{bne} \cdot \text{BrEQ}$

Controller: OR logic



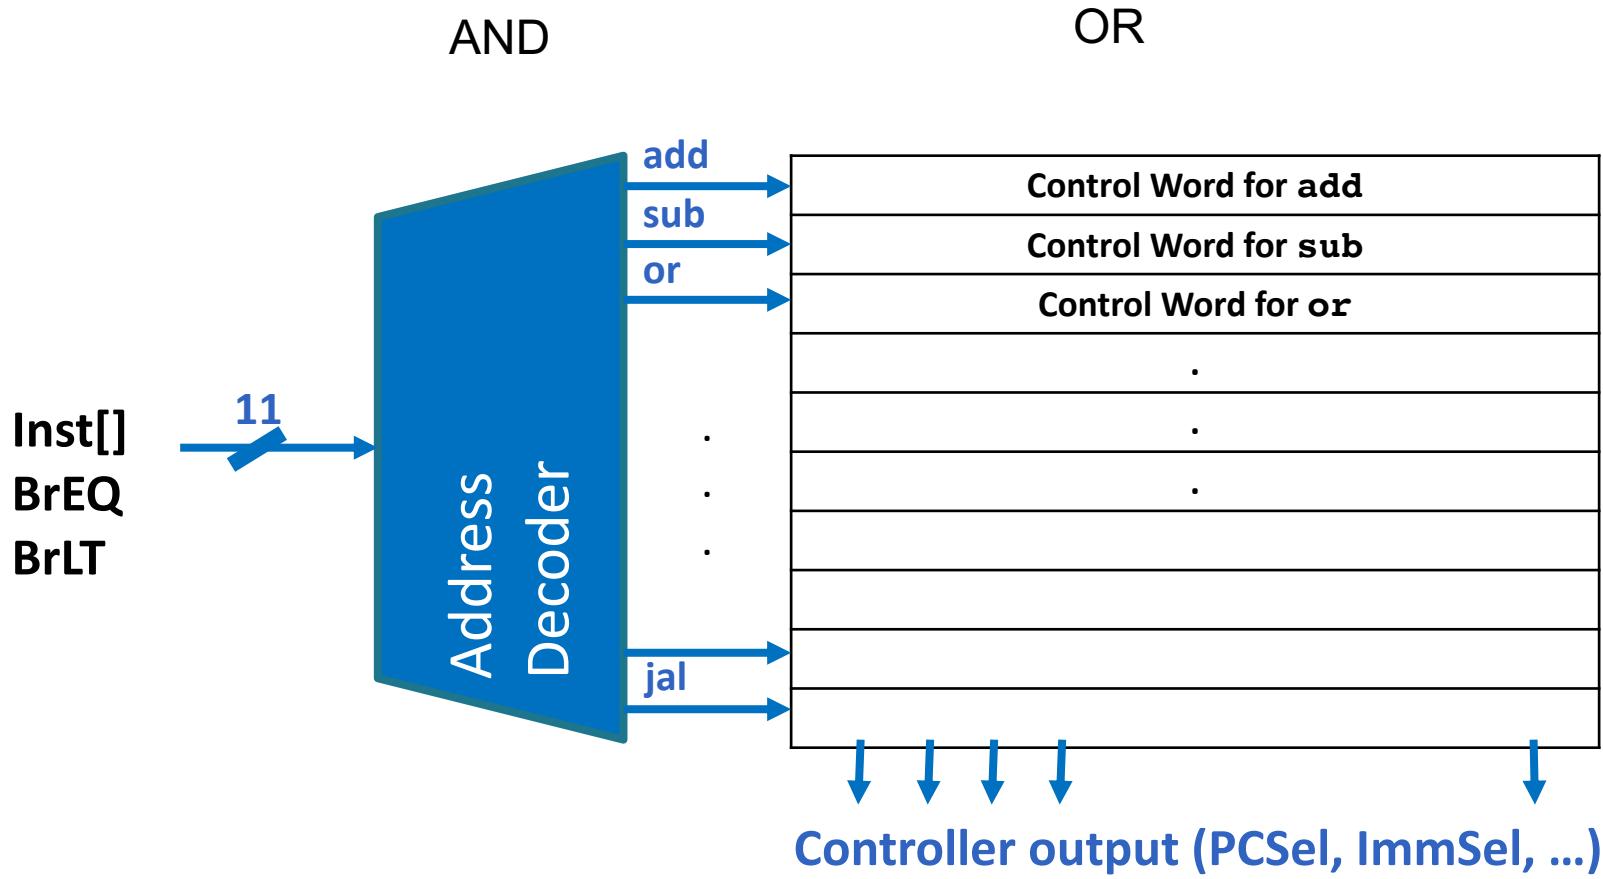
- MemRW = sw + sh + sb
- WBSel[0] = add + (all reg. to reg.) + AUIPC + ...
- WBSel[1] = JALR + JAL
- PCSel = beq_t + bne_t + blt_t + bltu_t + jal + jalr
- ...



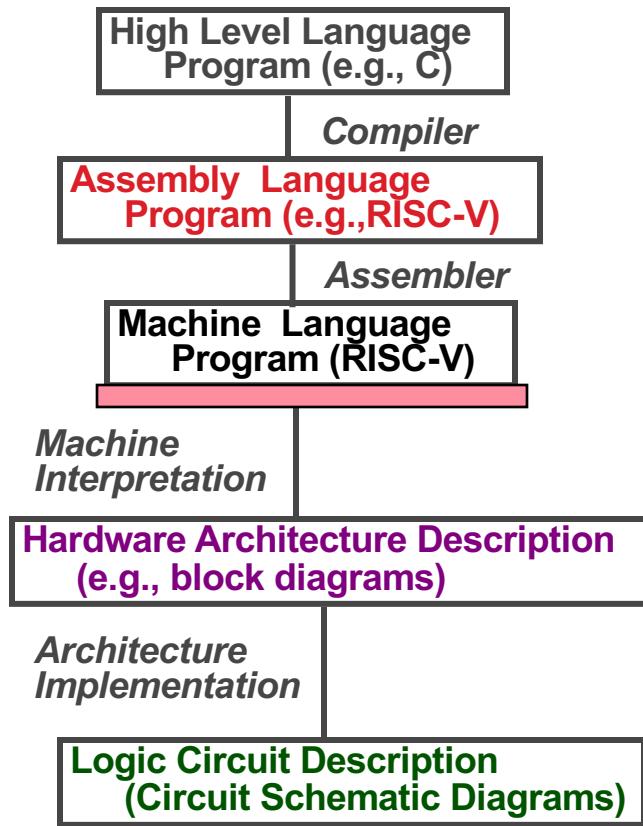
Control Logic Truth Table

Inst[31:0]	BrEq	BrLT	PCSel	ImmSel	BrUn	ASel	BSel	ALUSel	MemRW	RegWEn	WBSel
add	*	*	+4	*	*	Reg	Reg	Add	Read	1	ALU
sub	*	*	+4	*	*	Reg	Reg	Sub	Read	1	ALU
(R-R Op)	*	*	+4	*	*	Reg	Reg	(Op)	Read	1	ALU
addi	*	*	+4	I	*	Reg	Imm	Add	Read	1	ALU
lw	*	*	+4	I	*	Reg	Imm	Add	Read	1	Mem
sw	*	*	+4	S	*	Reg	Imm	Add	Write	0	*
beq	0	*	+4	B	*	PC	Imm	Add	Read	0	*
beq	1	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	0	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	1	*	+4	B	*	PC	Imm	Add	Read	0	*
blt	*	1	ALU	B	0	PC	Imm	Add	Read	0	*
bltu	*	1	ALU	B	1	PC	Imm	Add	Read	0	*
jalr	*	*	ALU	I	*	Reg	Imm	Add	Read	1	PC+4
jal	*	*	ALU	J	*	PC	Imm	Add	Read	1	PC+4
auiopc	*	*	+4	U	*	PC	Imm	Add	Read	1	ALU ⁵⁰

ROM Controller Implementation



Call home, we've made HW/SW contact!



“And In conclusion...”

- We have built a processor!
 - Capable of executing all RISC-V instructions in one cycle each
 - Not all units (hardware) used by all instructions
 - Critical path changes
- 5 Phases of execution
 - IF, ID, EX, MEM, WB
 - Not all instructions are active in all phases
- Controller specifies how to execute instructions
 - Implemented as ROM or logic